



# Linear Regression Continued



# Regularization



# Regularization

- Regularization seeks to solve a few common model issues by:
  - Minimizing model complexity
  - Penalizing the loss function
  - Reducing model overfitting (add more bias to reduce model variance)



# Regularization

- In general, we can think of regularization as a way to reduce model overfitting and variance.
  - Requires some additional bias
  - Requires a search for optimal penalty hyperparameter.



# Regularization

- Three main types of Regularization:
  - L1 Regularization
    - LASSO Regression
  - L2 Regularization
    - Ridge Regression
  - Combining L1 and L2
    - Elastic Net



# Regularization

- L1 regularization adds a penalty equal to the **absolute value** of the magnitude of coefficients.
  - Limits the size of the coefficients.
  - Can yield sparse models where some coefficients can become zero.



# Regularization

- L1 regularization adds a penalty equal to the **absolute value** of the magnitude of coefficients.

$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\beta_j| = \text{RSS} + \lambda \sum_{j=1}^p |\beta_j|$$



# Regularization

- L2 regularization adds a penalty equal to the **square** of the magnitude of coefficients.
  - All coefficients are shrunk by the same factor.
  - Does not necessarily eliminate coefficients.





# Regularization

- L2 regularization adds a penalty equal to the **square** of the magnitude of coefficients.

$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2 = \text{RSS} + \lambda \sum_{j=1}^p \beta_j^2$$



# Regularization

- Elastic net combines L1 and L2 with the addition of an alpha parameter deciding the ratio between them:

$$\frac{\sum_{i=1}^n (y_i - x_i^J \hat{\beta})^2}{2n} + \lambda \left( \frac{1 - \alpha}{2} \sum_{j=1}^m \hat{\beta}_j^2 + \alpha \sum_{j=1}^m |\hat{\beta}_j| \right)$$



# Regularization

- These regularization methods do have a cost:
  - Introduce an additional hyperparameter that needs to be tuned.
  - A multiplier to the penalty to decide the “strength” of the penalty.



# Regularization

- Later on, we will actually cover L2 regularization (Ridge Regression) first, due to the intuition behind the squared term being easier to understand.



# Regularization

- Before we dive straight into coding regularization with Scikit-Learn, we need to discuss a few more relevant topics:
  - Feature Scaling
  - Cross Validation



# Feature Scaling



# Feature Scaling

- Feature scaling provides many benefits to our machine learning process!
- Some machine learning models that rely on distance metrics (e.g. KNN) **require** scaling to perform well.
- Let's discuss the main ideas behind feature scaling...



# Feature Scaling

- Feature scaling improves the convergence of steepest descent algorithms, which do not possess the property of scale invariance.
- If features are on different scales, certain weights may update faster than others since the feature values  $x_j$  play a role in the weight updates.





# Feature Scaling

- Critical benefit of feature scaling related to gradient descent.
- There are some ML Algos where scaling won't have an effect (e.g. CART based methods).



# Feature Scaling

- Scaling the features so that their respective ranges are uniform is important in comparing measurements that have different units.
- Allows us directly compare model coefficients to each other.



# Feature Scaling

- Feature scaling caveats:
  - Must always scale new unseen data before feeding to model.
  - Effects direct interpretability of feature coefficients
    - Easier to compare coefficients to one another, harder to relate back to original unscaled feature.



# Feature Scaling

- Feature scaling benefits:
  - Can lead to great increases in performance.
  - Absolutely necessary for some models.
  - Virtually no “real” downside to scaling features.



# Feature Scaling

- Two main ways to scale features:
  - Standardization
    - Rescales data to have a mean ( $\mu$ ) of 0 and standard deviation ( $\sigma$ ) of 1.
  - Normalization
    - Rescales all data values to be between 0-1.



# Feature Scaling

- Standardization:
  - Rescales data to have a mean ( $\mu$ ) of 0 and standard deviation ( $\sigma$ ) of 1 (unit variance).

$$X_{changed} = \frac{X - \mu}{\sigma}$$



# Feature Scaling

- Standardization:
  - Namesake can be confusing since this is also referred to as “Z-score normalization”.

$$X_{changed} = \frac{X - \mu}{\sigma}$$



# Feature Scaling

- Normalization:
  - Scales all data values to be between 0 and 1.

$$X_{changed} = \frac{X - X_{min}}{X_{max} - X_{min}}$$





# Feature Scaling

- Normalization:
  - Simple and easy to understand.

$$X_{changed} = \frac{X - X_{min}}{X_{max} - X_{min}}$$



# Feature Scaling

- There are many more methods of scaling features and Scikit-Learn provides easy to use classes that “fit” and “transform” feature data for scaling.
- Let’s quickly discuss the fit and transform calls in more detail when it comes to scaling.



# Feature Scaling

- A `.fit()` method call simply calculates the necessary statistics ( $X_{min}$ ,  $X_{max}$ , mean, standard deviation).
- A `.transform()` call actually scales data and returns the new scaled version of data.
- Previously saw a similar process for polynomial feature conversion.



# Feature Scaling

- Very important consideration for fit and transform:
  - We only **fit** to training data.
  - Calculating statistical information should only come from training data.
  - Don't want to assume prior knowledge of the test set!



# Feature Scaling

- Using the full data set would cause **data leakage**:
  - Calculating statistics from full data leads to some information of the test set leaking into the training process upon transform() conversion.



# Feature Scaling

- Feature scaling process:
  - Perform train test split
  - **Fit to training** feature data
  - Transform training feature data
  - Transform test feature data



# Feature Scaling

- Do we need to scale the label( $y$ )?
  - In general it is not necessary nor advised.
  - Normalising the output distribution is altering the definition of the target.
  - Predicting a distribution that doesn't mirror your real-world target.



# Feature Scaling

- Do we need to scale the label?
  - Can negatively impact stochastic gradient descent.
- **[stats.stackexchange.com/questions/111467](https://stats.stackexchange.com/questions/111467)**





# Feature Scaling

- Now that we understand the benefits of feature scaling, let's move on to understanding the benefits of cross-validation!



# Cross Validation



# Cross Validation

- Cross validation is a more advanced set of methods for splitting data into training and testing sets.
- Cross Validation Relevant Reading:
  - Section 5.1 of ISLR



# Cross Validation

- We understand the intuition behind performing a train test split, we want to fairly evaluate our model's performance on unseen data.
- Unfortunately this means we are not able to tune hyperparameters to the **entire** dataset.



# Cross Validation

- Is there a way we can achieve the following:
  - Train on all the data
  - Evaluate on all the data
- While it sounds impossible, we can achieve this with cross validation!
- Let's have an overview of the concept...



# Cross Validation

- Imagine our data set:

X			y
Area m <sup>2</sup>	Bedrooms	Bathrooms	Price
200	3	2	\$500,000
190	2	1	\$450,000
230	3	3	\$650,000
180	1	1	\$400,000
210	2	2	\$550,000



# Cross Validation

- Let's convert this data into colored blocks for cross-validation

**X**

**y**

Area m <sup>2</sup>	Bedrooms	Bathrooms	Price
200	3	2	\$500,000
190	2	1	\$450,000
230	3	3	\$650,000
180	1	1	\$400,000
210	2	2	\$550,000



# Cross Validation

- Convert to generalized form

<b>x</b>			<b>y</b>
<b><math>x_1</math></b>	<b><math>x_2</math></b>	<b><math>x_3</math></b>	<b>y</b>
$x_1^1$	$x_1^1$	$x_1^1$	$y_1$
$x_1^2$	$x_1^2$	$x_1^2$	$y_2$
$x_1^3$	$x_1^3$	$x_1^3$	$y_3$
$x_1^4$	$x_1^4$	$x_1^4$	$y_4$
$x_1^5$	$x_1^5$	$x_1^5$	$y_5$





# Cross Validation

- Color based off train vs. test set.

	<b>x</b>			<b>y</b>
	<b>x<sub>1</sub></b>	<b>x<sub>2</sub></b>	<b>x<sub>3</sub></b>	<b>y</b>
<b>TRAIN</b>	$x_1^1$	$x_1^1$	$x_1^1$	$y_1$
	$x_1^2$	$x_1^2$	$x_1^2$	$y_2$
	$x_1^3$	$x_1^3$	$x_1^3$	$y_3$
<b>TEST</b>	$x_1^4$	$x_1^4$	$x_1^4$	$y_4$
	$x_1^5$	$x_1^5$	$x_1^5$	$y_5$



# Cross Validation

- Color based off train vs. test set.

	<b>x</b>			<b>y</b>
	<b>x<sub>1</sub></b>	<b>x<sub>2</sub></b>	<b>x<sub>3</sub></b>	<b>y</b>
<b>TRAIN</b>	$x_1^1$	$x_1^1$	$x_1^1$	$y_1$
	$x_1^2$	$x_1^2$	$x_1^2$	$y_2$
	$x_1^3$	$x_1^3$	$x_1^3$	$y_3$
<b>TEST</b>	$x_1^4$	$x_1^4$	$x_1^4$	$y_4$
	$x_1^5$	$x_1^5$	$x_1^5$	$y_5$



# Cross Validation

- For now just consider training vs testing:

TRAIN	$x_1^1$	$x_1^1$	$x_1^1$	$y_1$
	$x_1^2$	$x_1^2$	$x_1^2$	$y_2$
	$x_1^3$	$x_1^3$	$x_1^3$	$y_3$
TEST	$x_1^4$	$x_1^4$	$x_1^4$	$y_4$
	$x_1^5$	$x_1^5$	$x_1^5$	$y_5$



# Cross Validation

- Now we have all data, colored by training set versus test set.

TRAIN				
TEST				



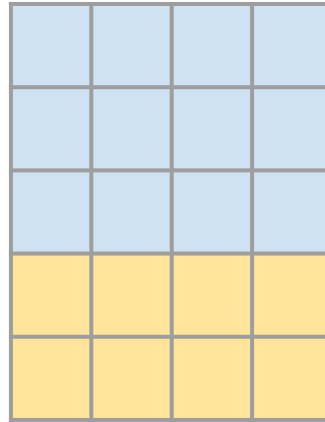
# Cross Validation

- Rotate and resize:




# Cross Validation

- Rotate and resize:





# Cross Validation

- Rotate and resize:





# Cross Validation

- Now we can represent full data and splits:







# Cross Validation

- Let's start with the entire original data:





# Cross Validation

- How does cross validation work?





# Cross Validation

- Split data into K equal parts:





# Cross Validation

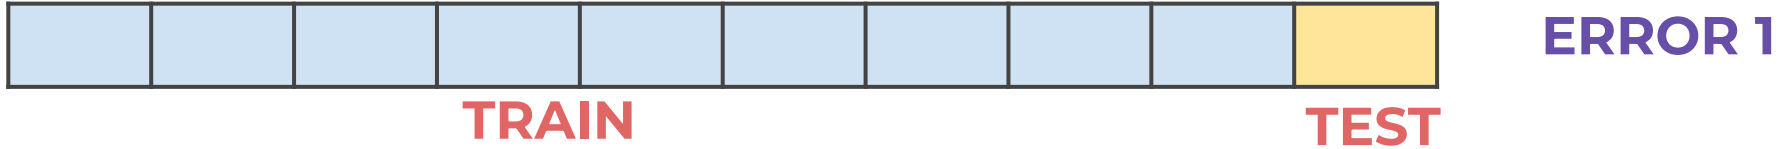
- $1/K$  left as test set





# Cross Validation

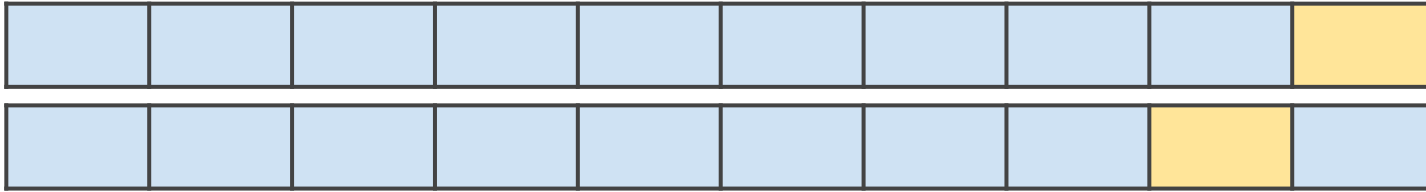
- Train model and get error metric for split:





# Cross Validation

- Repeat for another 1/K split



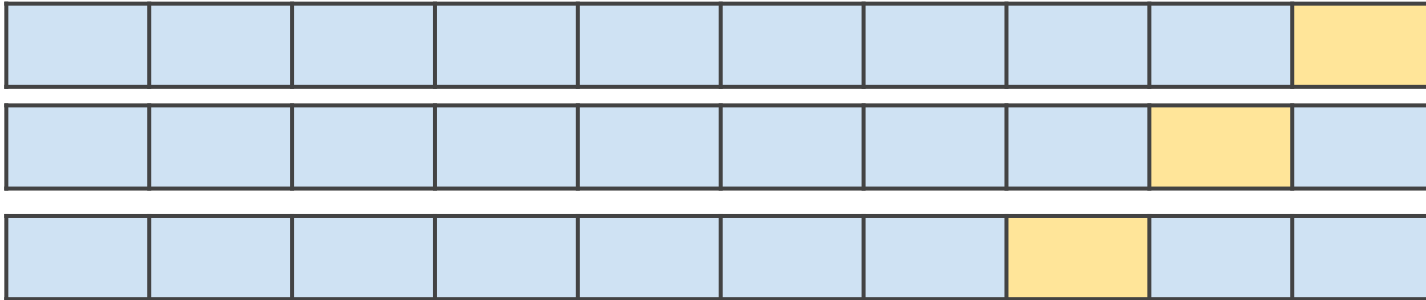
**ERROR 1**

**ERROR 2**



# Cross Validation

- Keep repeating for all possible splits



**ERROR 1**

**ERROR 2**

**ERROR 3**

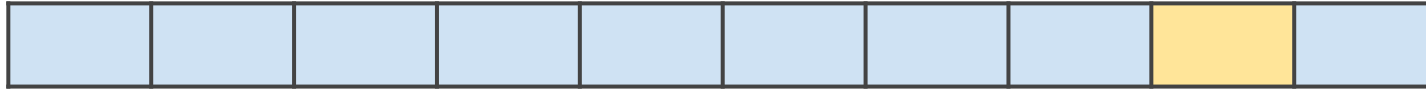


# Cross Validation

- Keep repeating for all possible splits



**ERROR 1**



**ERROR 2**



**ERROR 3**

...

...



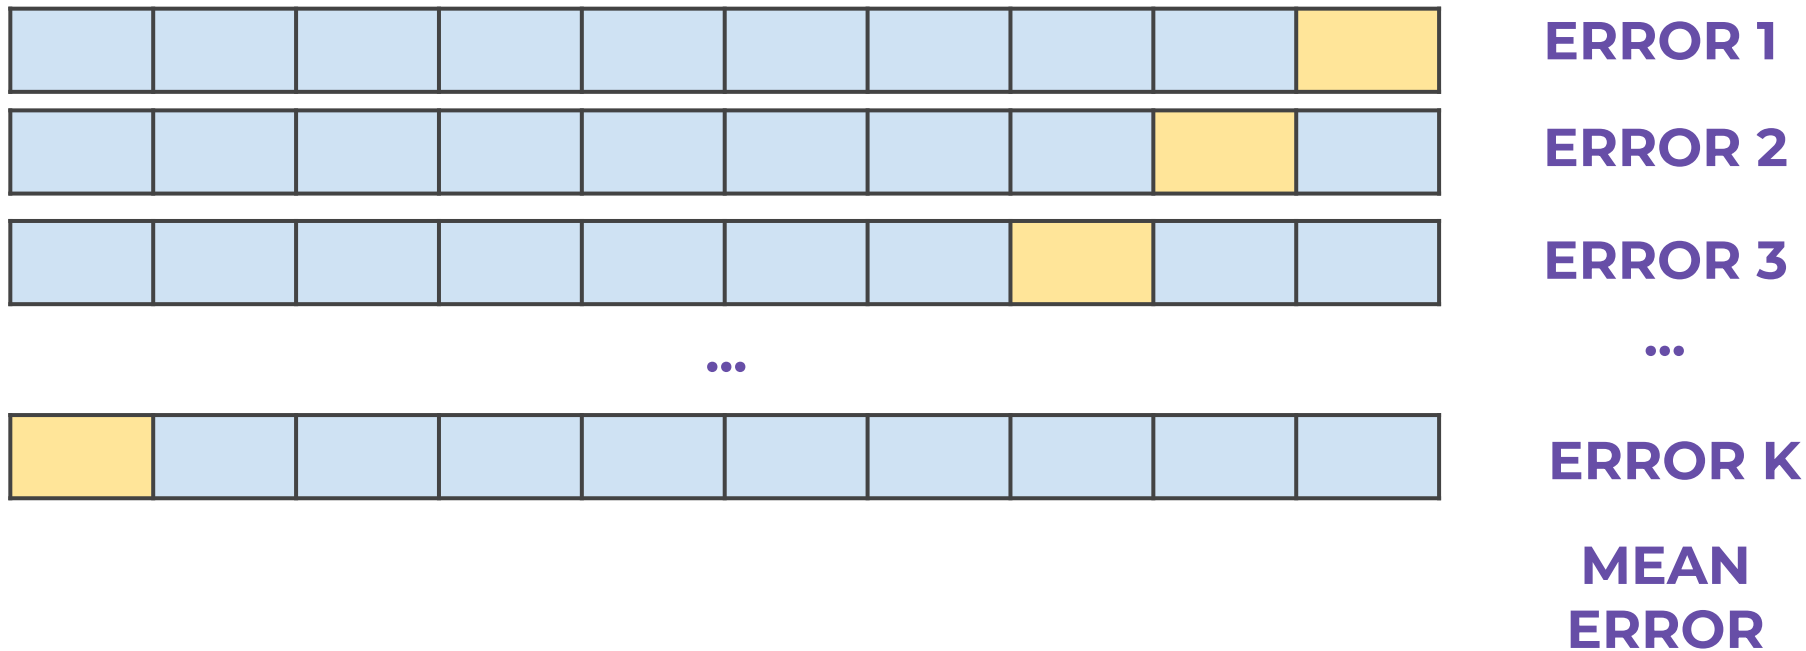
**ERROR K**





# Cross Validation

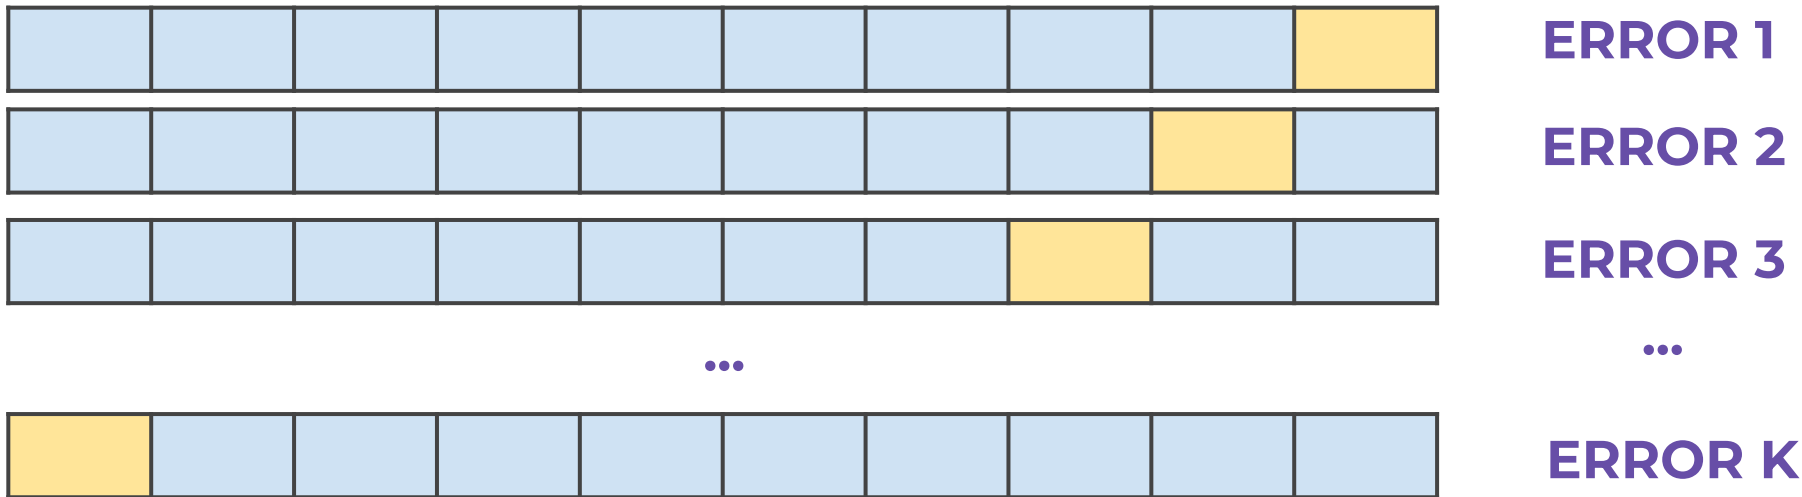
- Get average error





# Cross Validation

- Average error is the expected performance



MEAN



# Cross Validation

- We were able to train on all data **and** evaluate on all data!
- We get a better sense of true performance across multiple potential splits.
- What is the cost of this?
  - We have to repeat computations  $K$  number of times!



# Cross Validation

- This is known as K-fold cross-validation.
- Common choice for K is 10 so each test set is 10% of your total data.
- Largest K possible would be K equal to the number of number of rows.
  - This is known as **leave one out** cross validation.
  - Computationally expensive!



# Validation Set



# Cross Validation

- One consideration to note with K-fold cross validation and a standard train test split is *fairly tuning hyperparameters*.
- If we tune hyperparameters to test data performance, are we ever fairly getting performance metrics?



# Cross Validation

- How can we understand how the model behaves for data that it has not seen **and** not been influenced by for hyperparameter tuning?
- For this we can use a **hold out** test set.
- Let's explore what this looks like...



# Cross Validation

- Start with entire data set:







# Cross Validation

- Remove a hold out test set





# Cross Validation

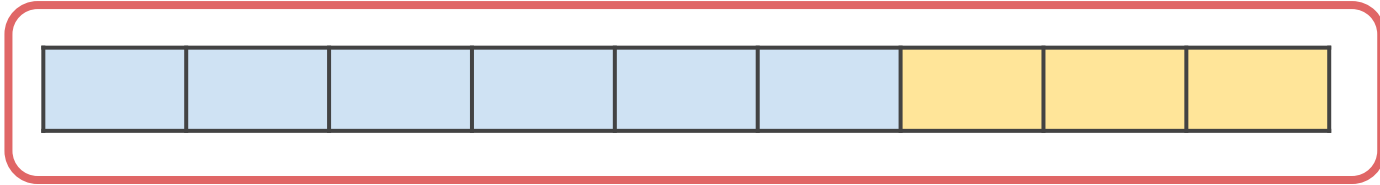
- Perform “classic” train test split:





# Cross Validation

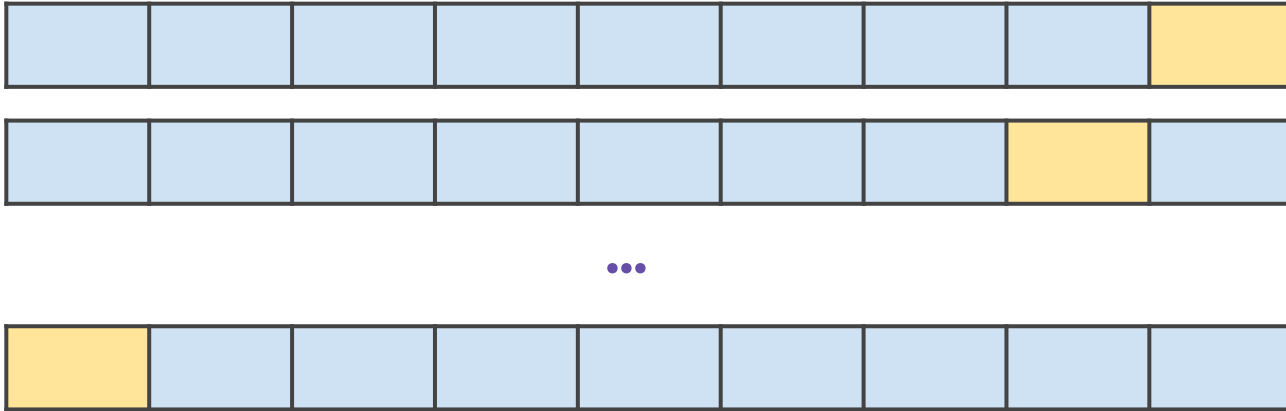
- Train and tune on this data:





# Cross Validation

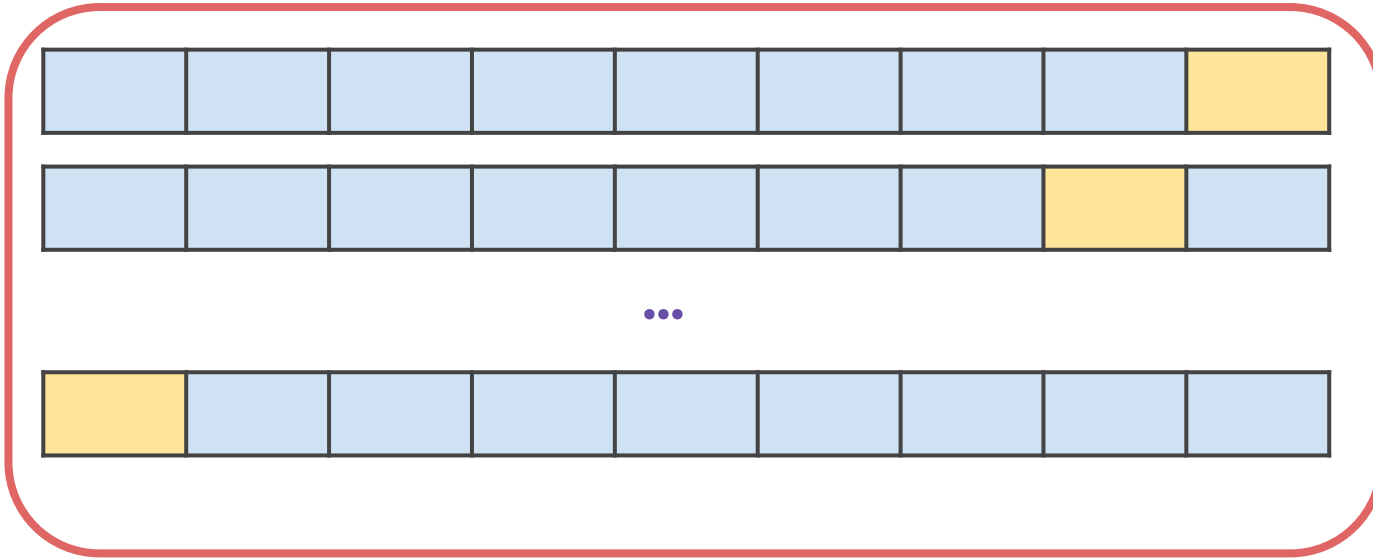
- Or K-Fold cross validation





# Cross Validation

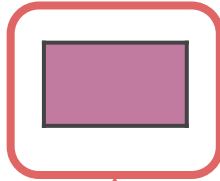
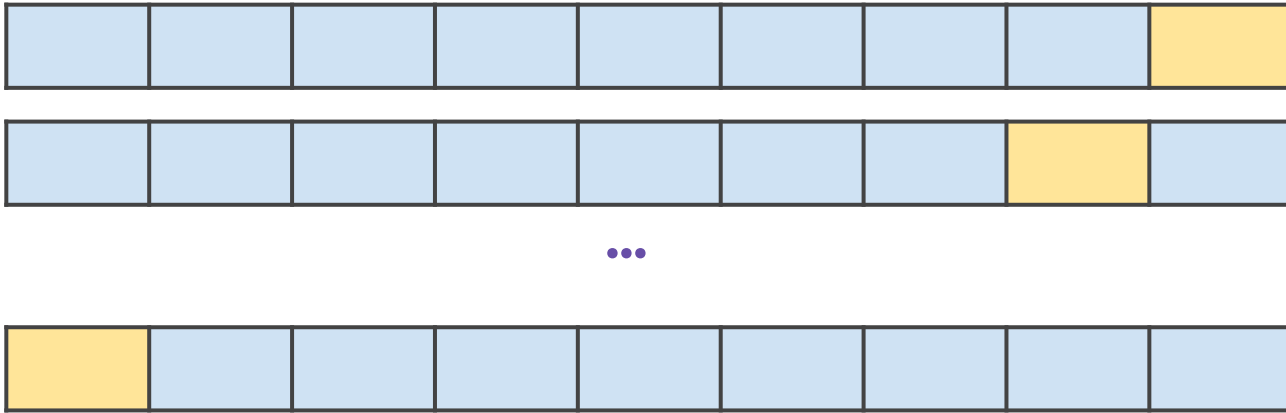
- Train **and** tune on this data:





# Cross Validation

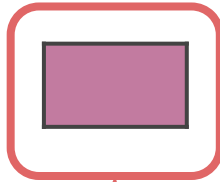
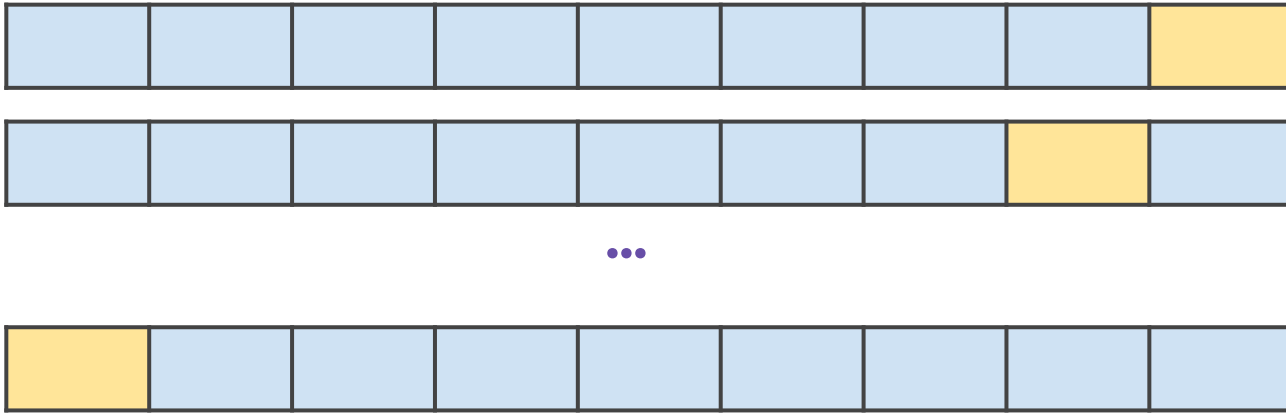
- **After** training and tuning perform **final evaluation** hold out test set.





# Cross Validation

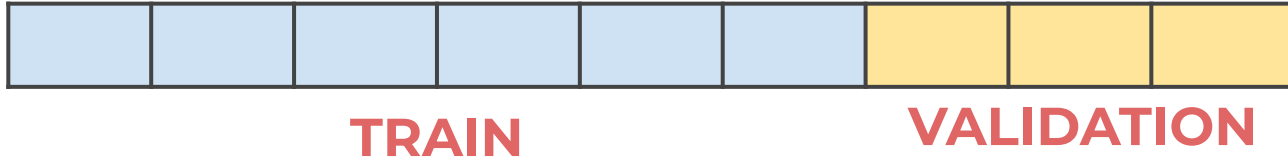
- Can **not** tune after this **final** test evaluation!





# Cross Validation

- Train | Validation | Test Split



- Allows us to get a true final performance metric to report.
- No editing model after this!





# Cross Validation

- All these approaches are valid, each situation is unique!
- Keep in mind:
  - Previous modeling work
  - Reporting requirements
  - Fairness of evaluation
  - Context of data and model



# Cross Validation

- Many regularization methods have tunable parameters we can adjust based on cross-validation techniques.
- For simplicity, there are times in the course we will opt for a simple two part train test split.



# Ridge Regression

Theory and Intuition



# Ridge Regression

- Ridge Regression is a regularization technique that *works by helping reduce the potential for overfitting to the training data.*
- It does this by adding in a penalty term to the error that is based on the squared value of the coefficients.



# Ridge Regression

- Ridge Regression is a regularization method for Linear Regression.
- Relevant Reading in ISLR:
  - Section 6.2.1
- Let's explore the main concepts behind how Ridge Regression works...



# Ridge Regression

- Recall the general formula for the regression line:

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x_1 + \hat{\beta}_2 x_2 + \cdots + \hat{\beta}_p x_p$$



# Ridge Regression

- These Beta coefficients were solved by minimizing the residual sum of squares (RSS).

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x_1 + \hat{\beta}_2 x_2 + \cdots + \hat{\beta}_p x_p$$



# Ridge Regression

- These Beta coefficients were solved by minimizing the residual sum of squares (RSS).

$$\text{RSS} = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$





# Ridge Regression

- We could substitute our regression equation for  $\hat{\mathbf{y}}$ :

$$\text{RSS} = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$



# Ridge Regression

- We could substitute our regression equation for  $\hat{\mathbf{y}}$ :

$$\begin{aligned}\text{RSS} &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 \\ &= \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_{i1} - \hat{\beta}_2 x_{i2} - \cdots - \hat{\beta}_p x_{ip})^2\end{aligned}$$



# Ridge Regression

- We can then summarize RSS as:

$$\text{RSS} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2$$



# Ridge Regression

- The goal of Ridge Regression is to help prevent overfitting by adding an additional penalty term.

$$\text{RSS} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2$$



# Ridge Regression

- Ridge Regression adds a **shrinkage penalty**:

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

- Ridge Regression seeks to minimize this entire error term **RSS + Penalty**.

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

- **Shrinkage penalty** based off the squared coefficient:

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

- **Shrinkage penalty** has a **tunable lambda parameter!**

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \boxed{\lambda} \sum_{j=1}^p \beta_j^2$$





# Ridge Regression

- Lambda determines how severe the penalty is.

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

- In theory it can be any value from 0 to positive infinity.

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

- If it is zero, then it is simply back to RSS.

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

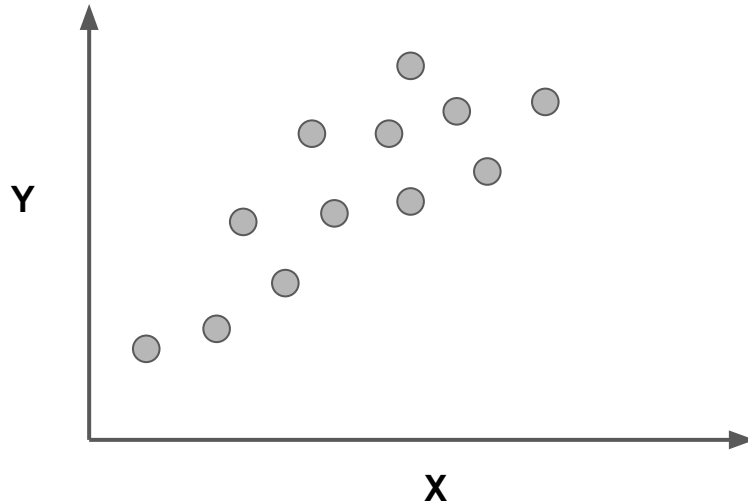
- Let's explore a simple thought experiment to get an intuition behind Ridge Regression...

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

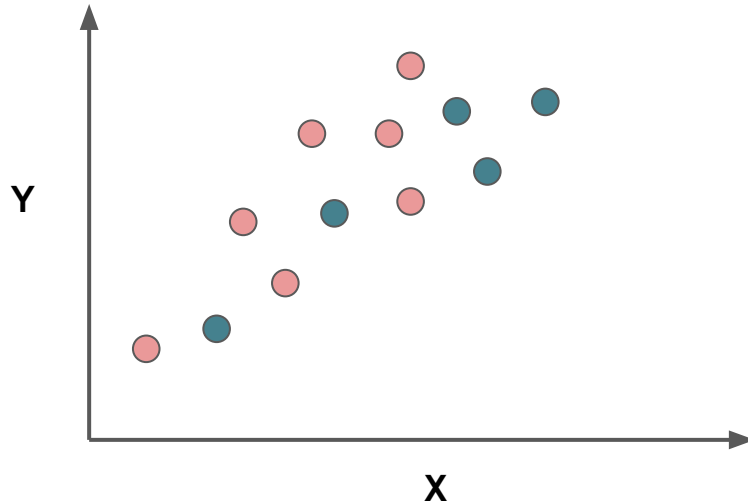
- Imagine the following data set.





# Ridge Regression

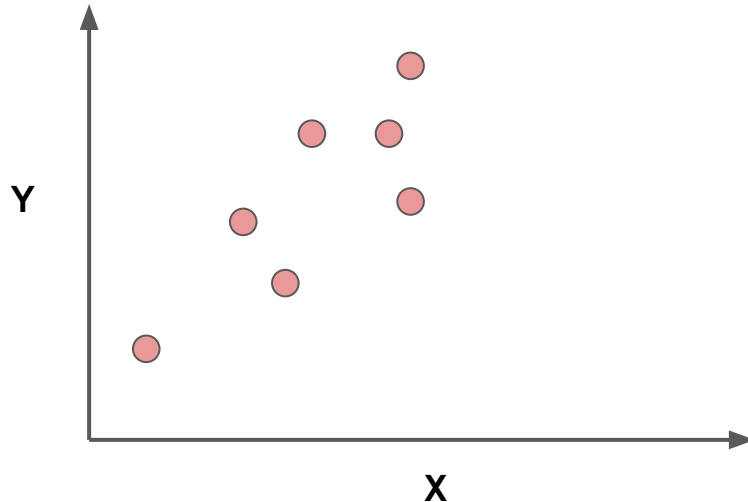
- We can split it into a training set and test set:





# Ridge Regression

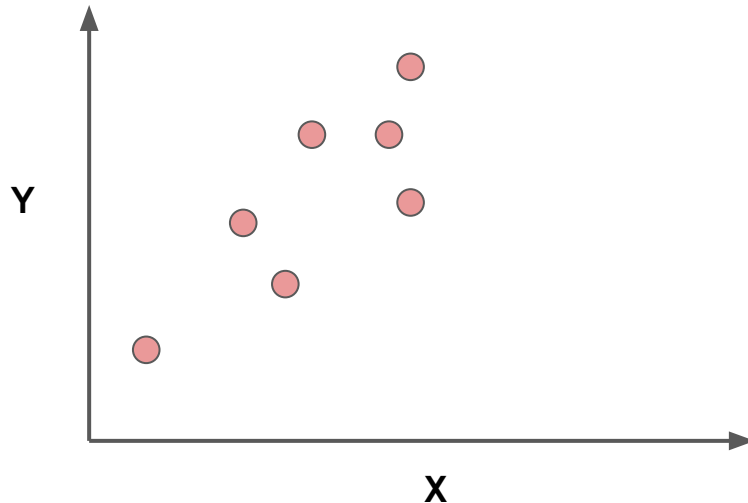
- Now we can fit on the training data to produce the line:  $\hat{y} = \beta_1 x + \beta_0$





# Ridge Regression

- Regardless of RSS or Ridge error, we're still trying to create a line:  $\hat{y} = \beta_1 x + \beta_0$

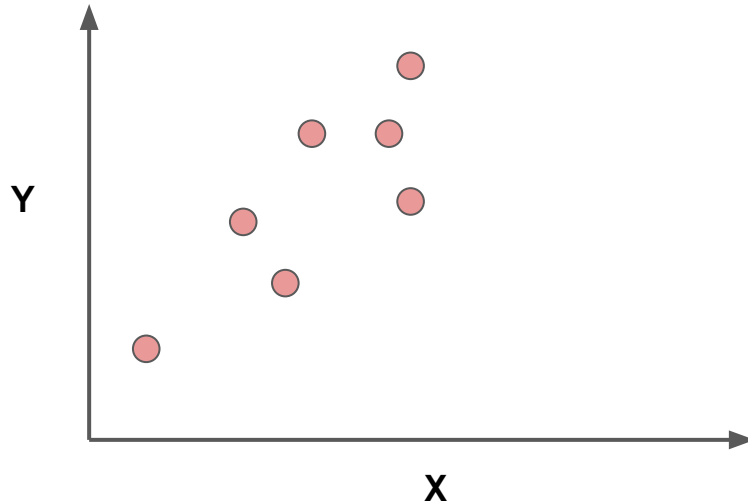






# Ridge Regression

- The only difference would be the coefficients found.

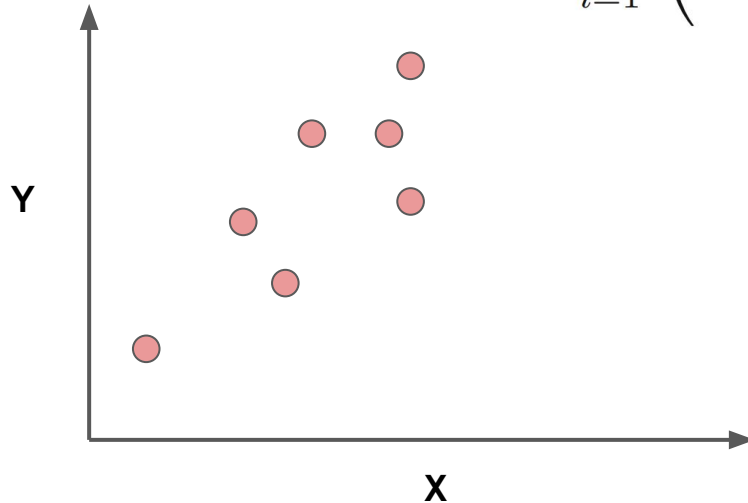




# Ridge Regression

- First let's fit using only RSS...

$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2$$

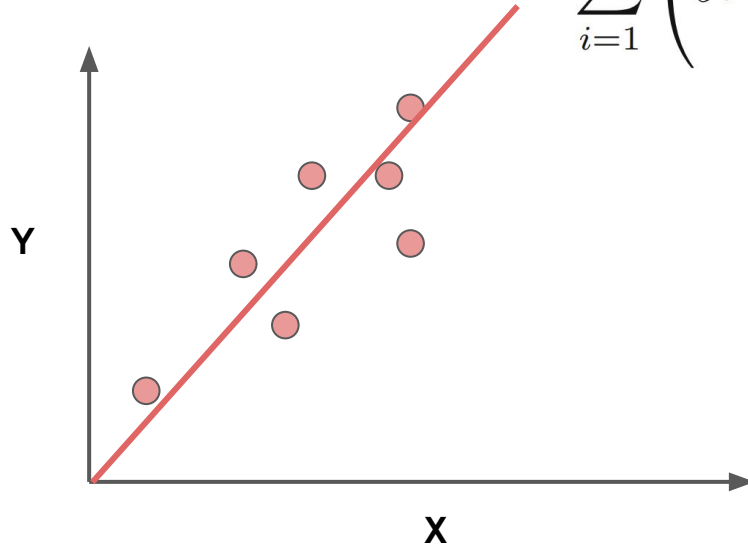




# Ridge Regression

- Our fitted  $\hat{y} = \beta_1 x + \beta_0$

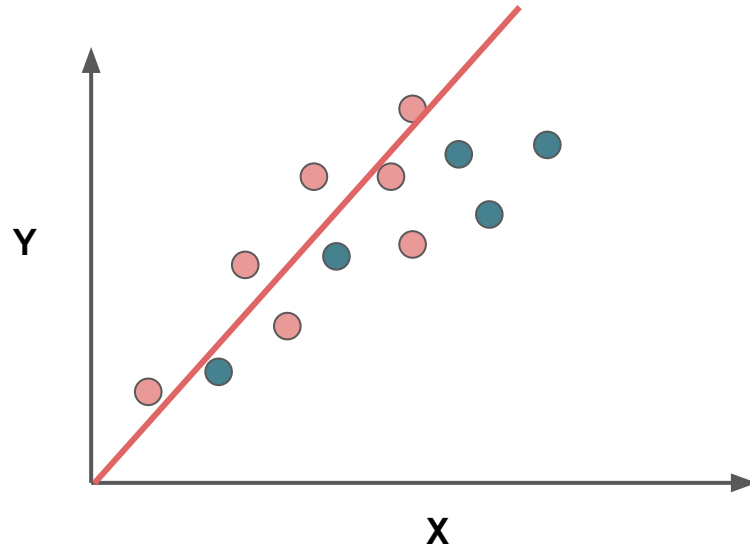
$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2$$





# Ridge Regression

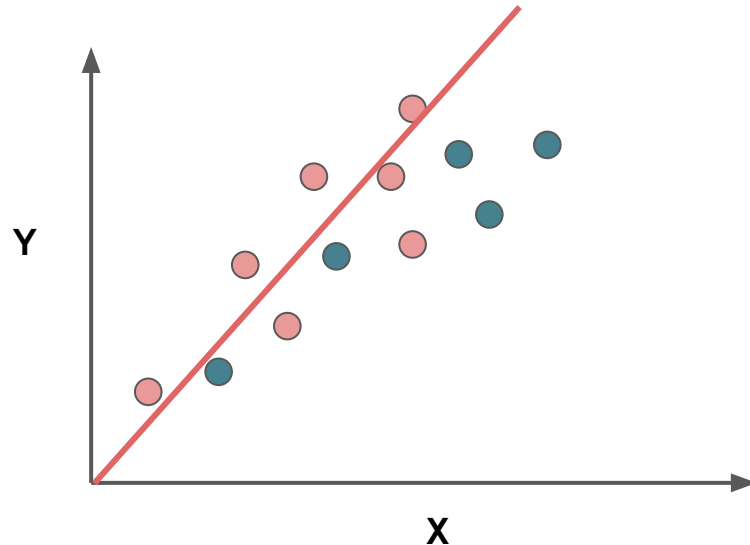
- Appears to have over fit to training data.





# Ridge Regression

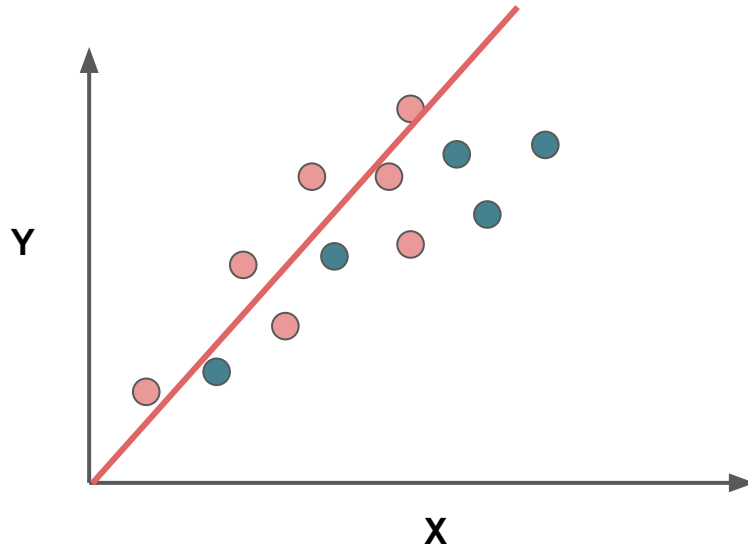
- This means we have high **variance**.





# Ridge Regression

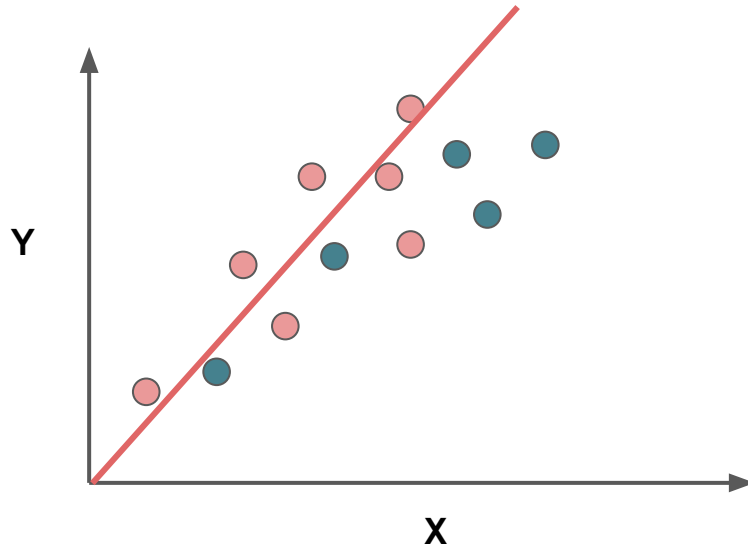
- We know there is a **bias-variance** trade-off.





# Ridge Regression

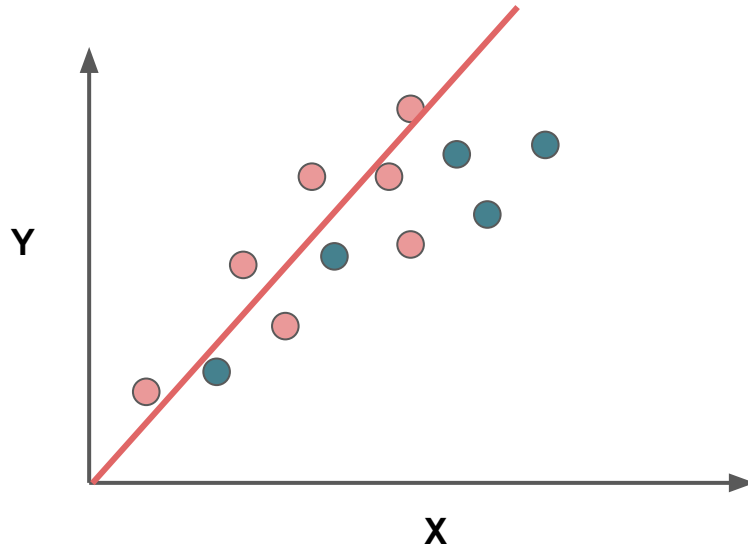
- But could we introduce a little more **bias** to significantly **reduce** variance?





# Ridge Regression

- Would adding the penalty term help generalize with more **bias**?

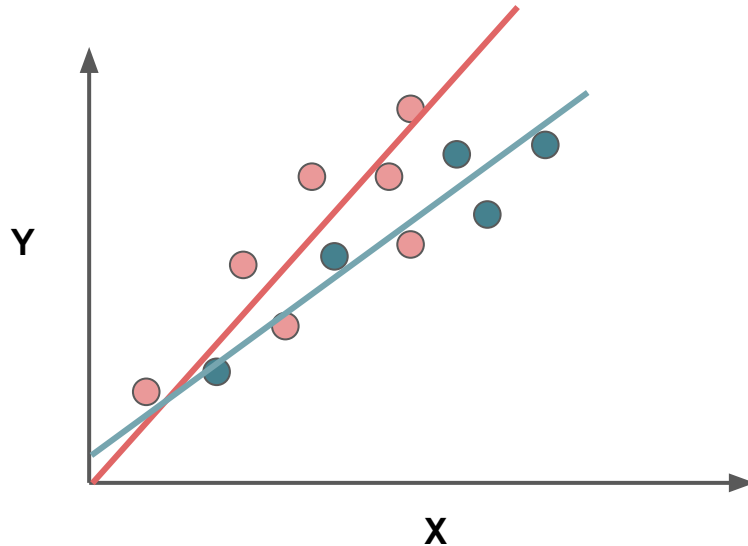






# Ridge Regression

- Adding bias can help generalize  $\hat{y} = \beta_1 x + \beta_0$

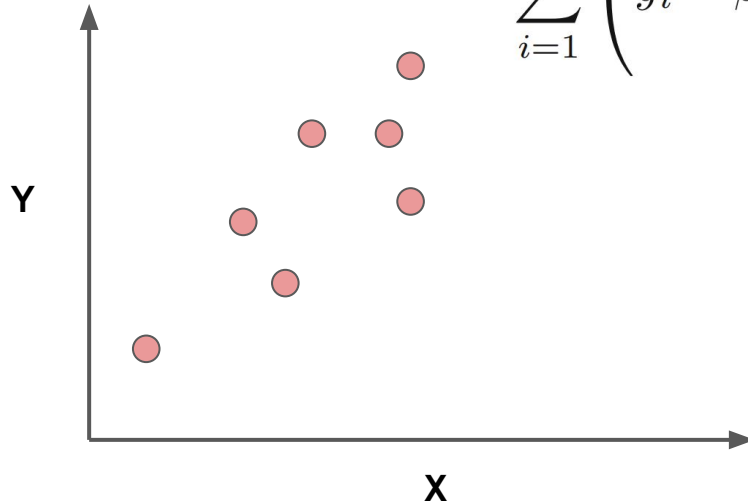




# Ridge Regression

- Let's imagine trying to reduce the Ridge Regression error term:

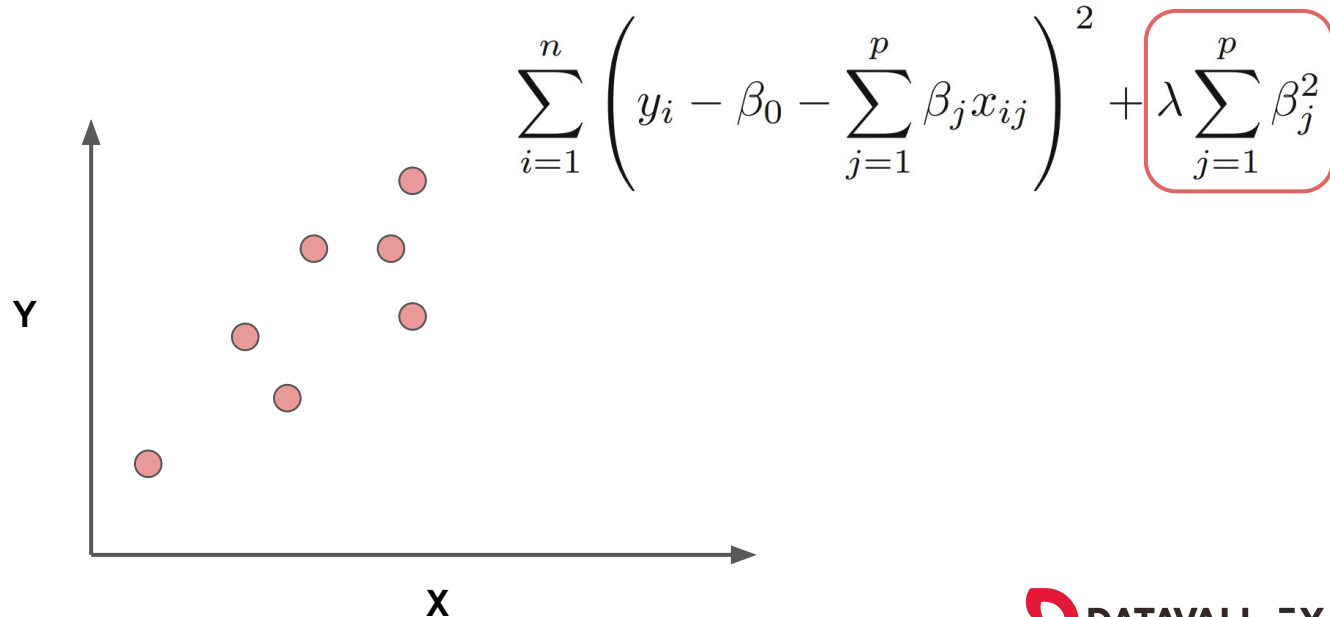
$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$





# Ridge Regression

- There is  $\lambda$  and the squared slope coefficient.

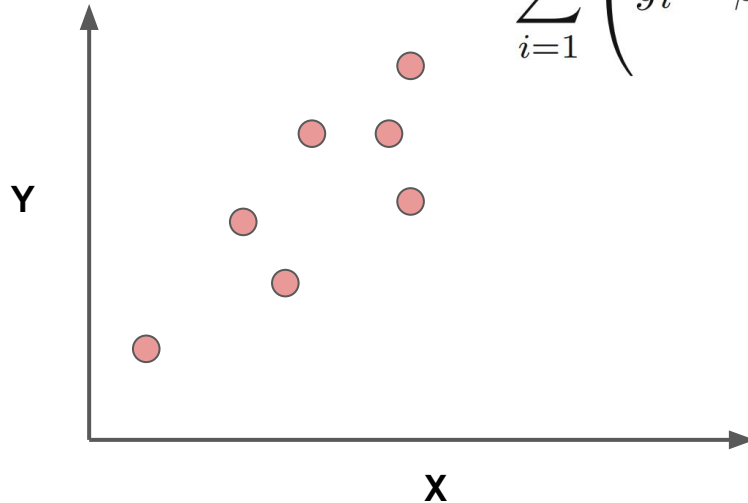




# Ridge Regression

- In the case of  $\hat{y} = \beta_1 x + \beta_0$

$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$

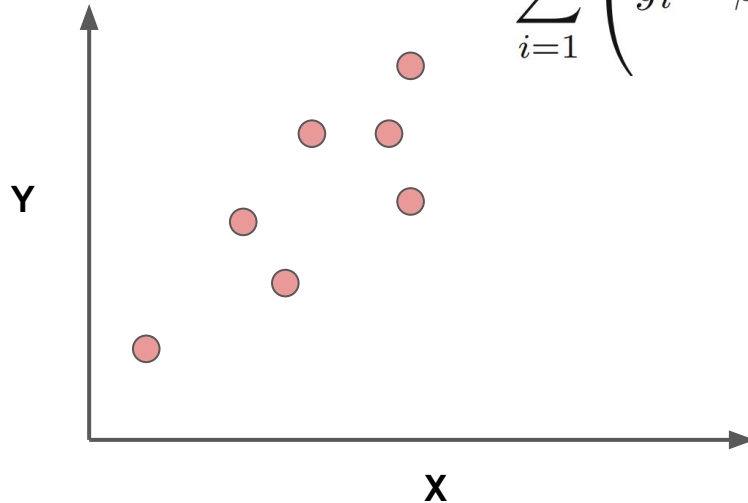




# Ridge Regression

- In the case of  $\hat{y} = \beta_1 x + \beta_0$

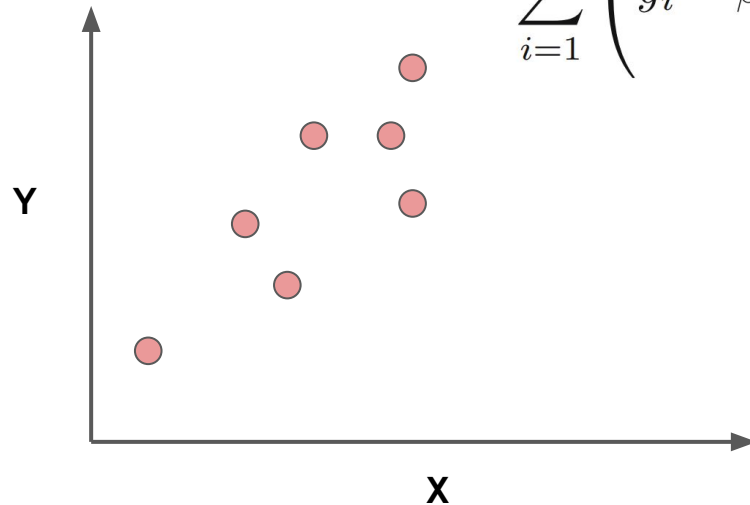
$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$





# Ridge Regression

- Let's assume  $\lambda = 1$



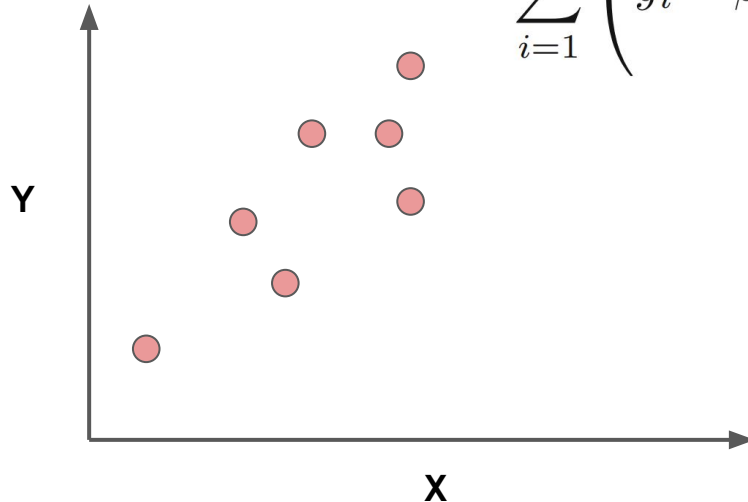
$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

- This punishes a large slope for  $\hat{y} = \beta_1 x + \beta_0$

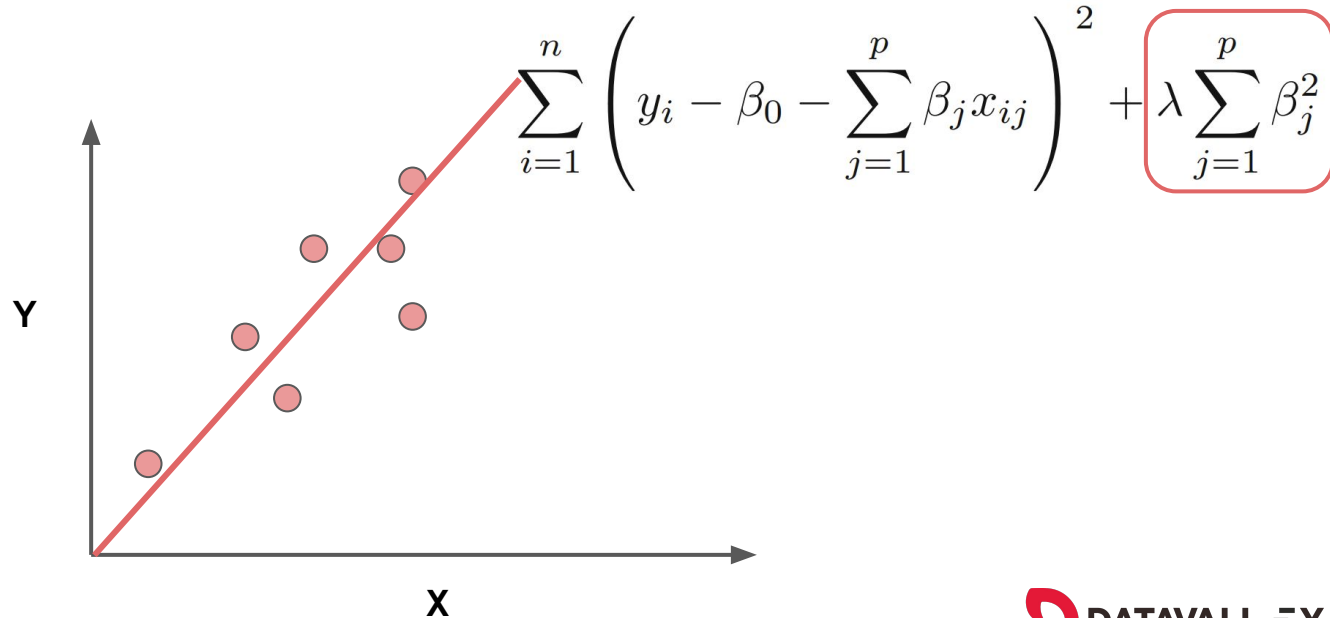
$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$





# Ridge Regression

- For single feature this lowers slope

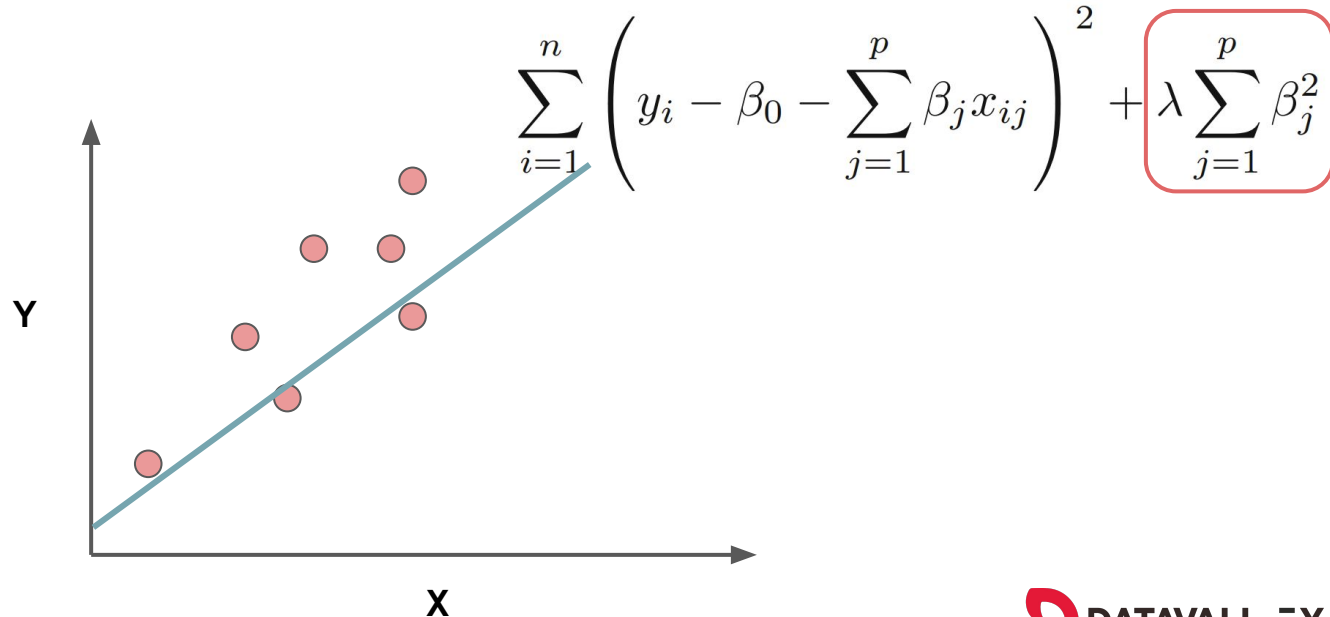






# Ridge Regression

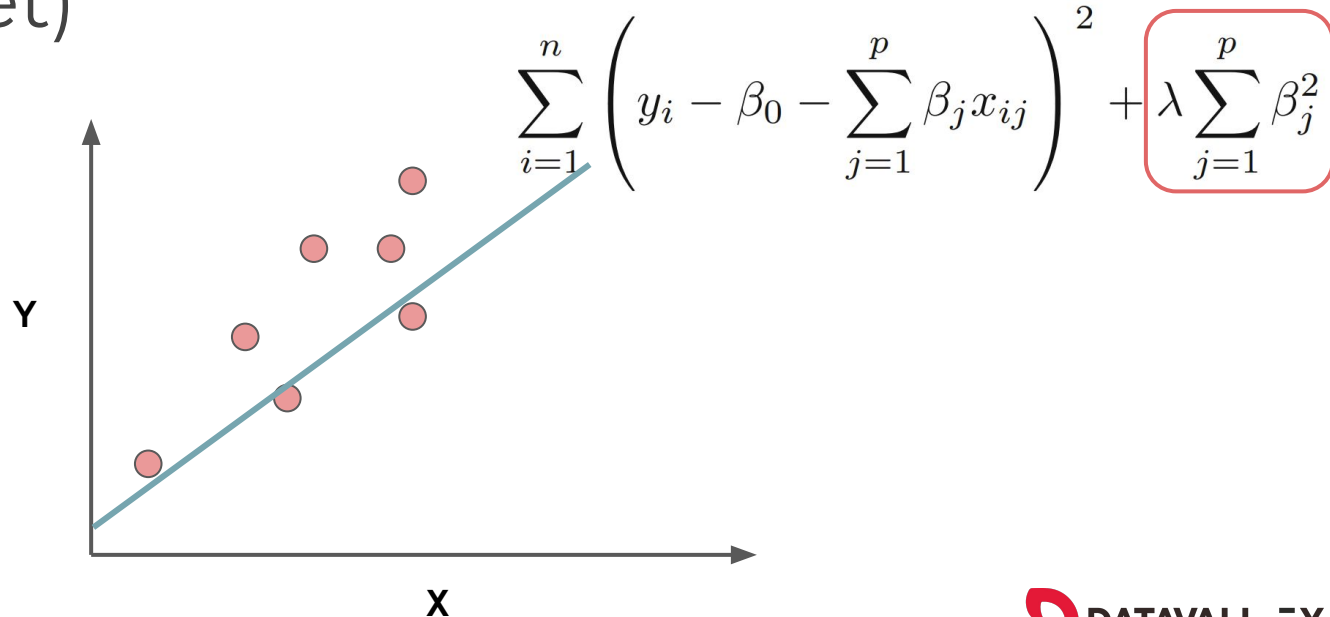
- For single feature this lowers slope





# Ridge Regression

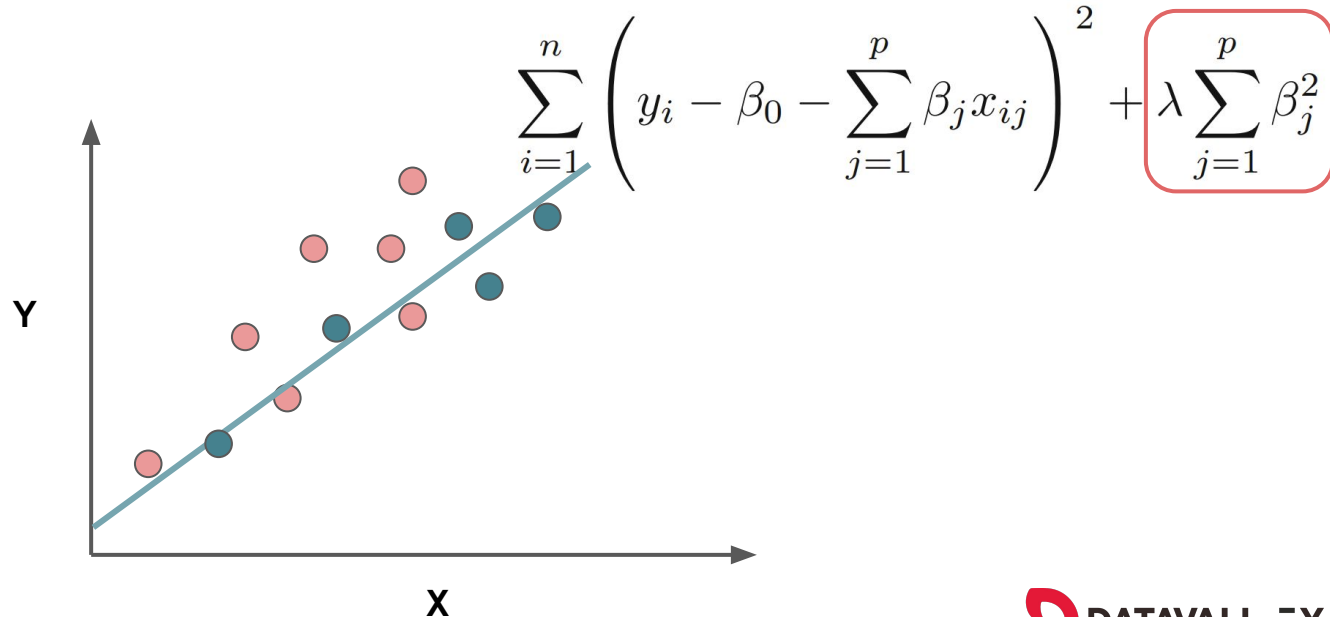
- At the cost of some additional bias (error in training set)





# Ridge Regression

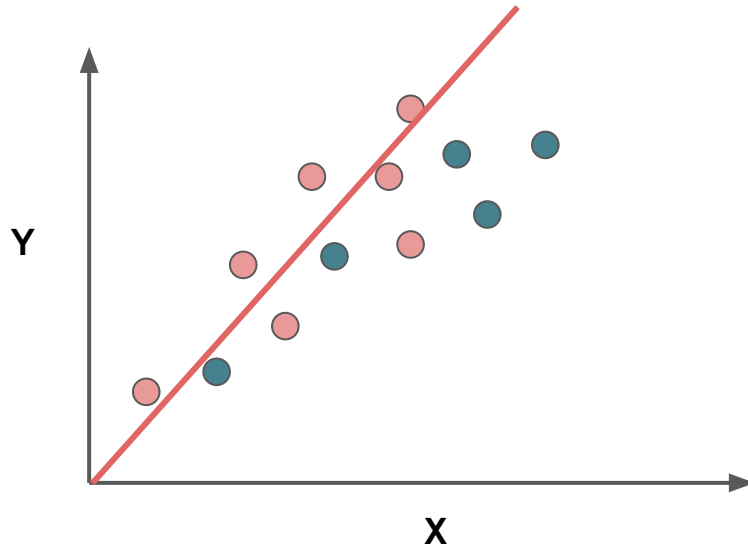
- We generalize better to unseen data





# Ridge Regression

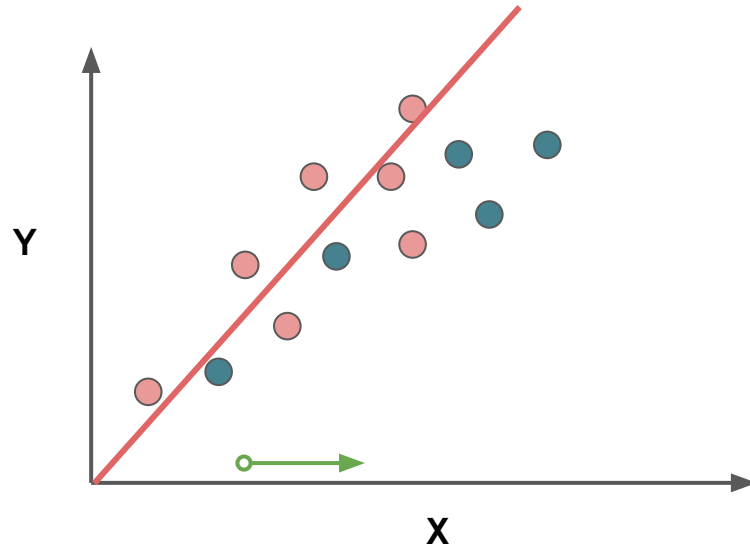
- Consider overfitting to training set:





# Ridge Regression

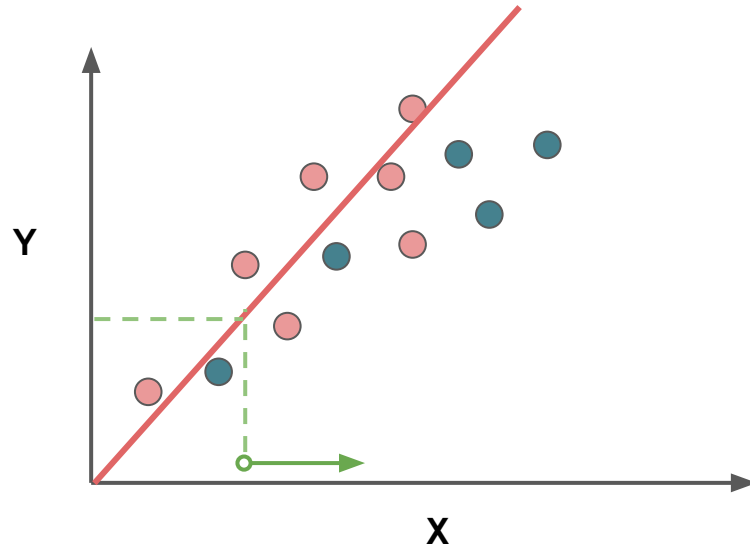
- An increase in  $X$  results in a greater  $y$  response:





# Ridge Regression

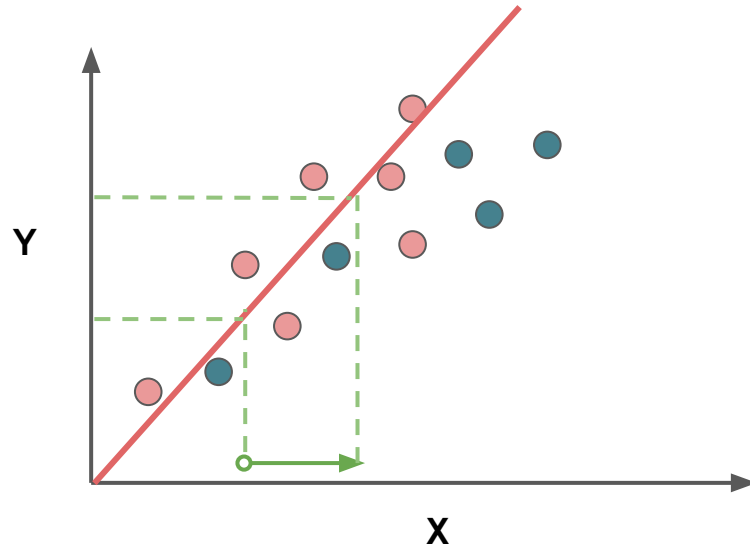
- An increase in  $X$  results in a greater  $y$  response:





# Ridge Regression

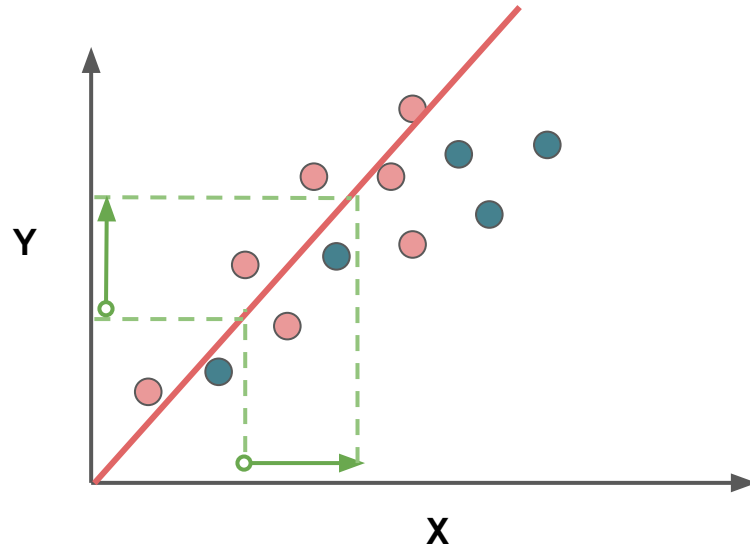
- An increase in  $X$  results in a greater  $y$  response:





# Ridge Regression

- An increase in  $X$  results in a greater  $y$  response:

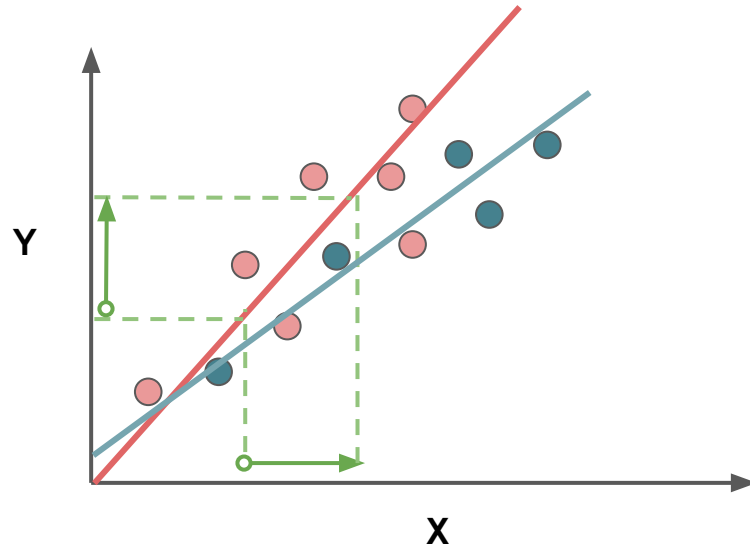






# Ridge Regression

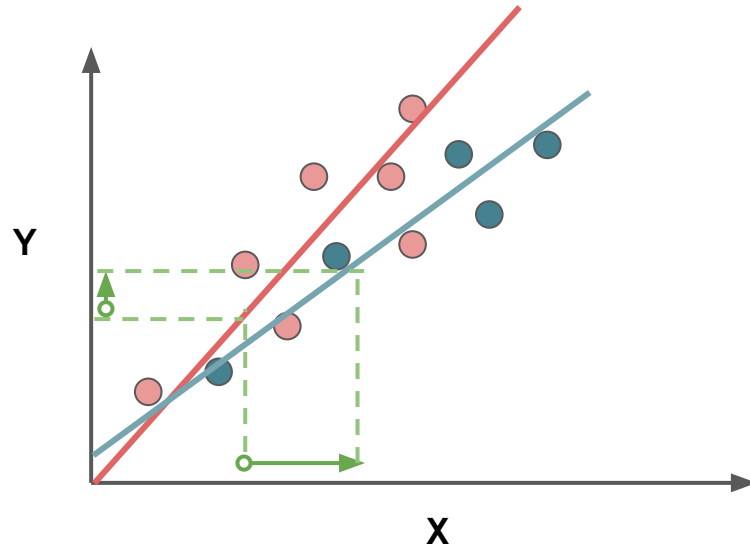
- Compare to a more generalized model that used Ridge Regression:





# Ridge Regression

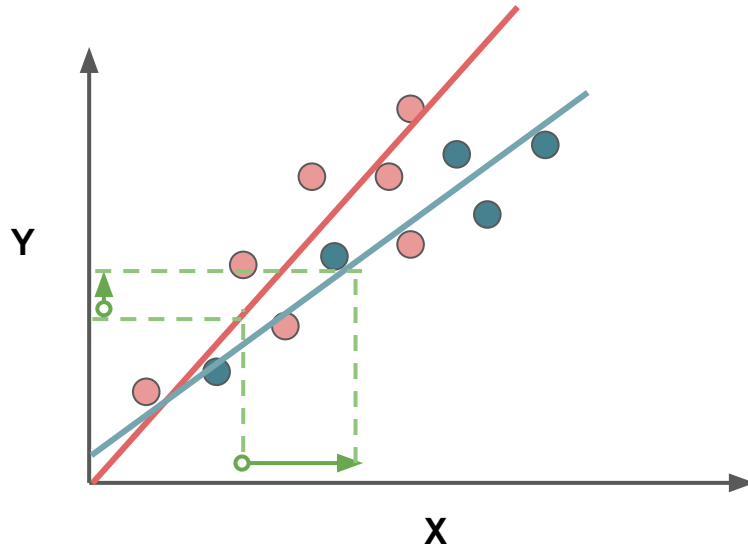
- Same feature change does not produce as much y response:





# Ridge Regression

- Trying to minimize a squared Beta term leads us to punish larger coefficients.

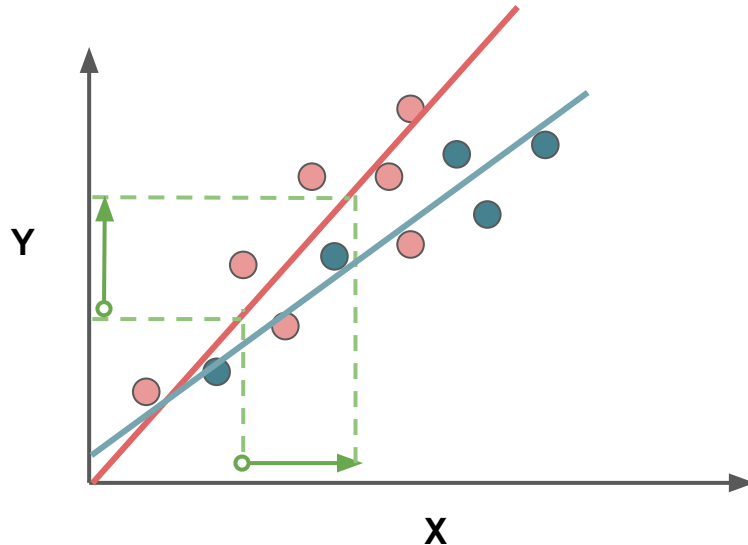


$$\lambda \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

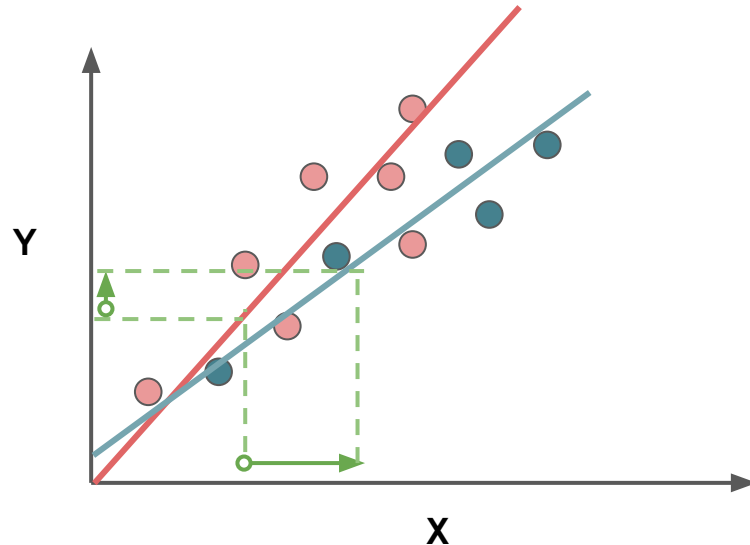
- In the case of a single feature, a larger Beta means a steeper sloped line.





# Ridge Regression

- A steeper sloped line would mean more response per increase in X value.

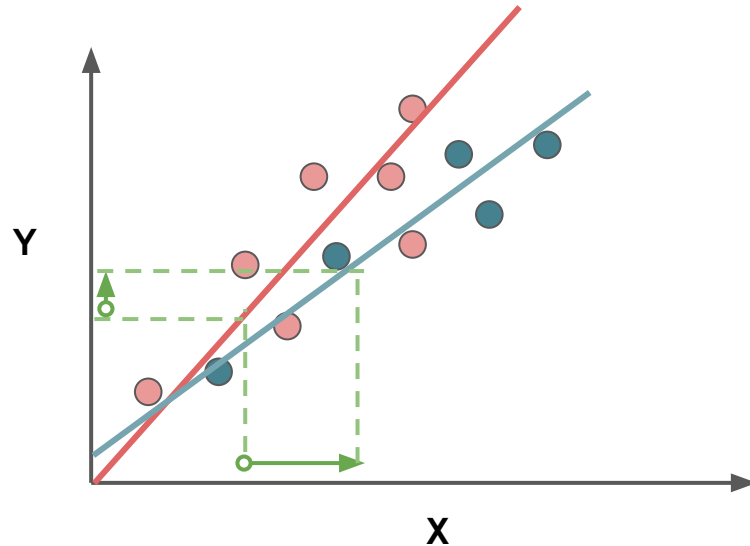


$$\lambda \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

- What about the lambda term? How much should we punish these larger coefficients?



$$\lambda \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

- We simply use cross-validation to explore multiple lambda options and then choose the best one!

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

L2 Regularization with Scikit-learn





# Ridge Regression

- Important Note!
  - Sklearn refers to lambda as alpha within the class call!

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \boxed{\lambda} \sum_{j=1}^p \beta_j^2$$



# Ridge Regression

- Important Note!
  - For cross validation metrics, sklearn uses a “scorer object”.
  - All scorer objects follow the convention that **higher** return values are **better** than lower return values.



# Ridge Regression

- Important Note!
  - For example, obviously higher accuracy is better.
  - But higher RMSE is actually worse!
  - So Scikit-Learn fixes this by using a **negative** RMSE as its scorer metric.



# Ridge Regression

- Important Note!
  - This allows for uniformity across **all** scorer metrics, even across different tasks types.
  - The same idea of uniformity across model classes applies to referring to the penalty strength parameter as **alpha**.



# Lasso Regression

L1 Regularization



# Regularization

- L1 regularization adds a penalty equal to the **absolute value** of the magnitude of coefficients.

$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\beta_j| = \text{RSS} + \lambda \sum_{j=1}^p |\beta_j|$$



# Regularization

- L1 regularization adds a penalty equal to the **absolute value** of the magnitude of coefficients.
  - Limits the size of the coefficients.
  - Can yield sparse models where some coefficients can become zero.



# Regularization

- LASSO can force some of the coefficient estimates to be exactly equal to zero when the tuning parameter  $\lambda$  is sufficiently large.
- Similar to subset selection, the LASSO performs variable selection.
- Models generated from the LASSO are generally much easier to interpret.





# Regularization

- LassoCV with sklearn operates on checking a number of alphas within a range, instead of providing the alphas directly.
- Let's explore the results of LASSO in Python and Scikit-Learn!



# Elastic Net

L1 and L2 Regularization



# Regularization

- We've been able to perform Ridge and Lasso regression.
- We know Lasso is able to shrink coefficients to zero, but we haven't taken a deeper dive into how or why that is.
- This ability becomes more clear when learning about **elastic net** which combines Lasso and Ridge together!



# Regularization

- We can rewrite Lasso and Ridge:

$$\underset{\beta}{\text{minimize}} \left\{ \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^p |\beta_j| \leq s$$

and

$$\underset{\beta}{\text{minimize}} \left\{ \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^p \beta_j^2 \leq s,$$



# Regularization

- There is some sum **s** which allows to rewrite the penalty as a requirement:

$$\underset{\beta}{\text{minimize}} \left\{ \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^p |\beta_j| \leq s$$

and

$$\underset{\beta}{\text{minimize}} \left\{ \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^p \beta_j^2 \leq s,$$



# Regularization

- There is some sum **s** which allows to rewrite the penalty as a requirement:

$$\underset{\beta}{\text{minimize}} \left\{ \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^p |\beta_j| \leq s$$

and

$$\underset{\beta}{\text{minimize}} \left\{ \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^p \beta_j^2 \leq s,$$





# Elastic Net

- Start with a simple thought experiment:
  - A simple equation:
    - $\hat{y} = \beta_1 x_1 + \beta_2 x_2$
  - We know that regularization can be expressed as an additional requirement that RSS is subject to.



# Elastic Net

- Start with a simple thought experiment:
  - A simple equation:
    - $\hat{y} = \beta_1 x_1 + \beta_2 x_2$
  - L1 constrains the sum of absolute values.
    - $\sum |\beta|$
  - L2 constrains the sum of squared values.
    - $\sum \beta^2$





# Elastic Net

- Start with a simple thought experiment:
  - A simple equation:
    - $\hat{y} = \beta_1 x_1 + \beta_2 x_2$
  - There is some sum **s** that the penalty is less than.



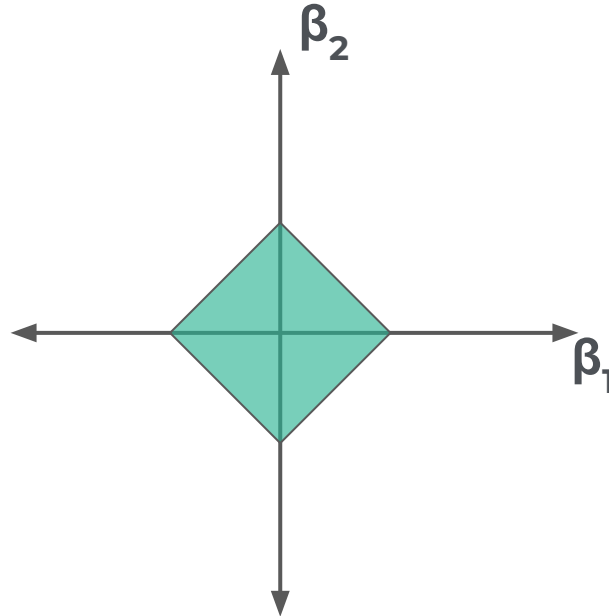
# Elastic Net

- For the case of only two features:
  - $\hat{y} = \beta_1 x_1 + \beta_2 x_2$
- Lasso Regression Penalty:
  - $|\beta_1| + |\beta_2| \leq s$
- Ridge Regression Penalty:
  - $\beta_1^2 + \beta_2^2 \leq s$



# Elastic Net

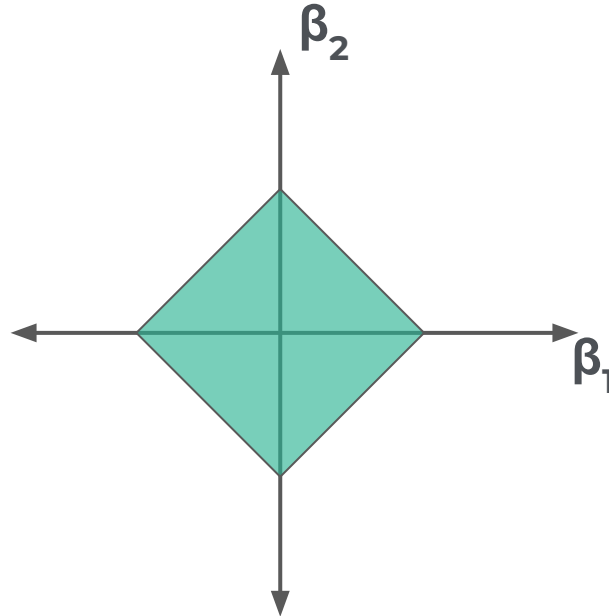
- Let's plot Lasso:  $|\beta_1| + |\beta_2| \leq s$





# Elastic Net

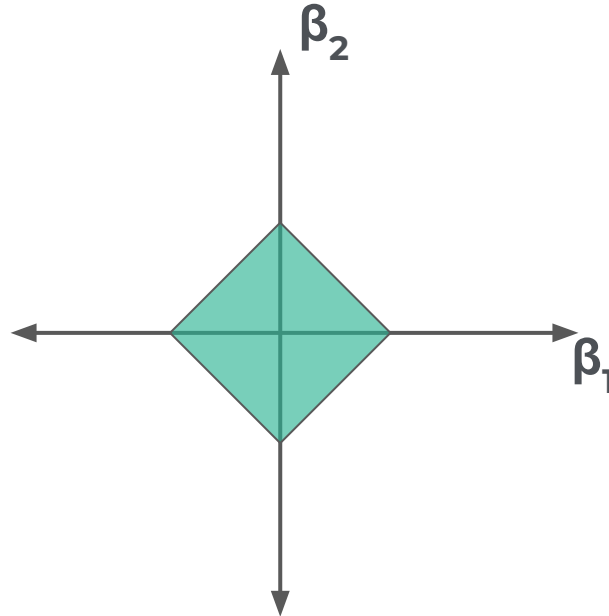
- Let's plot Lasso:  $|\beta_1| + |\beta_2| \leq s$





# Elastic Net

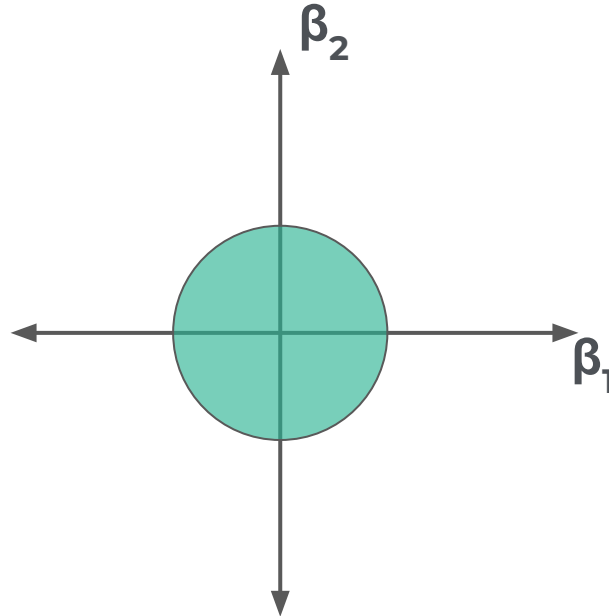
- Let's plot Lasso:  $|\beta_1| + |\beta_2| \leq s$





# Elastic Net

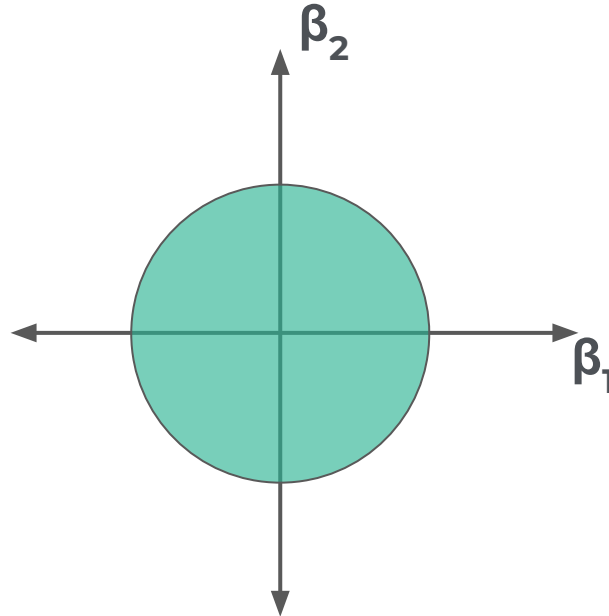
- Let's plot Ridge:  $\beta_1^2 + \beta_2^2 \leq s$





# Elastic Net

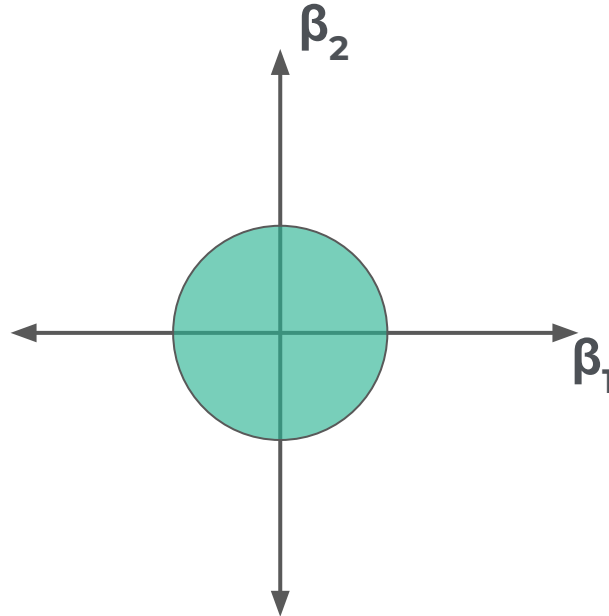
- Let's plot Ridge:  $\beta_1^2 + \beta_2^2 \leq s$





# Elastic Net

- Let's plot Ridge:  $\beta_1^2 + \beta_2^2 \leq s$

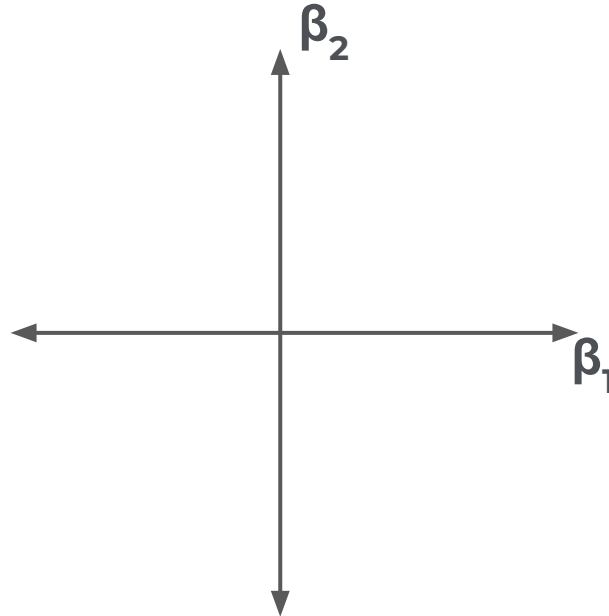






# Elastic Net

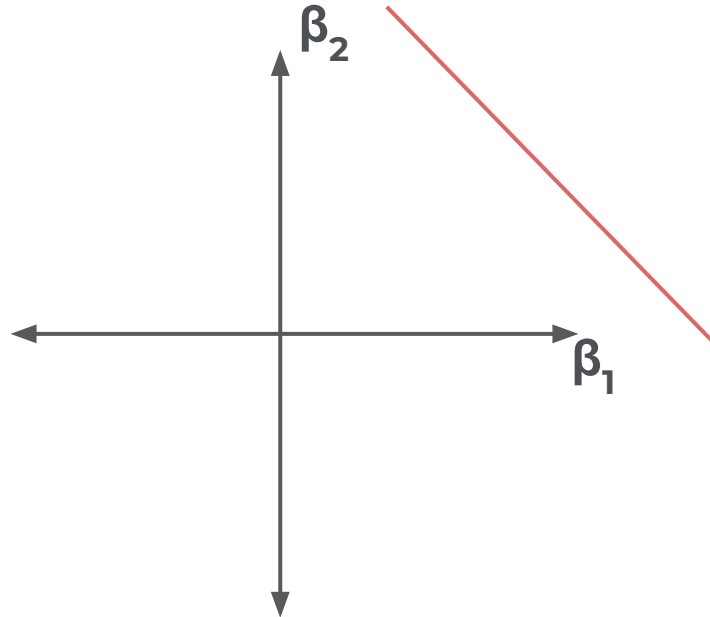
- What would RSS look like?





# Elastic Net

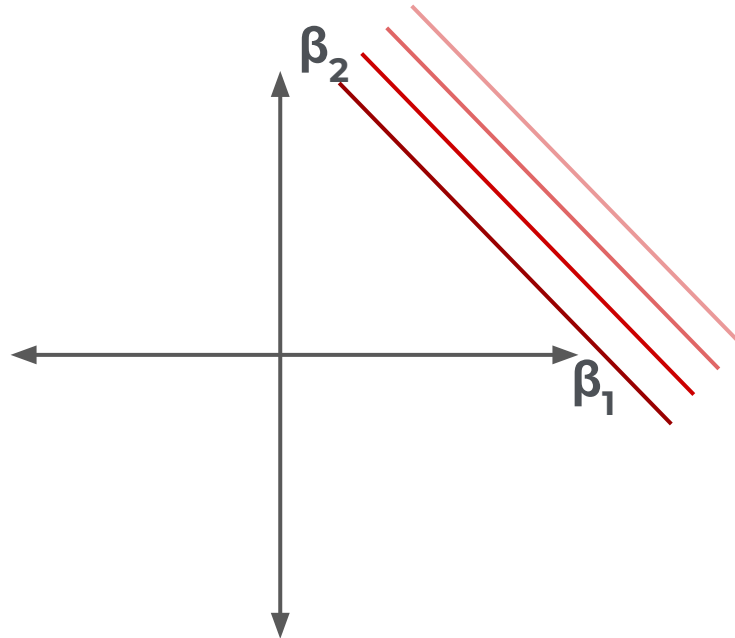
- What would RSS look like?





# Elastic Net

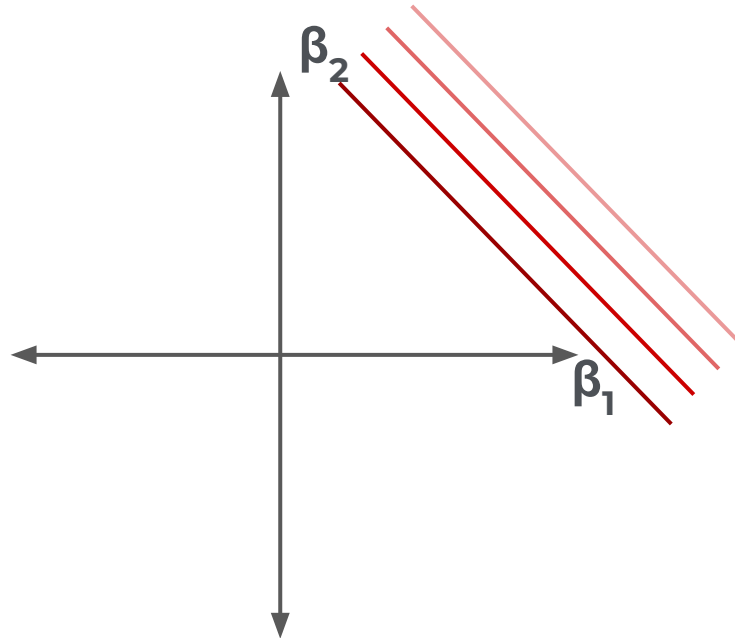
- What would RSS look like?





# Elastic Net

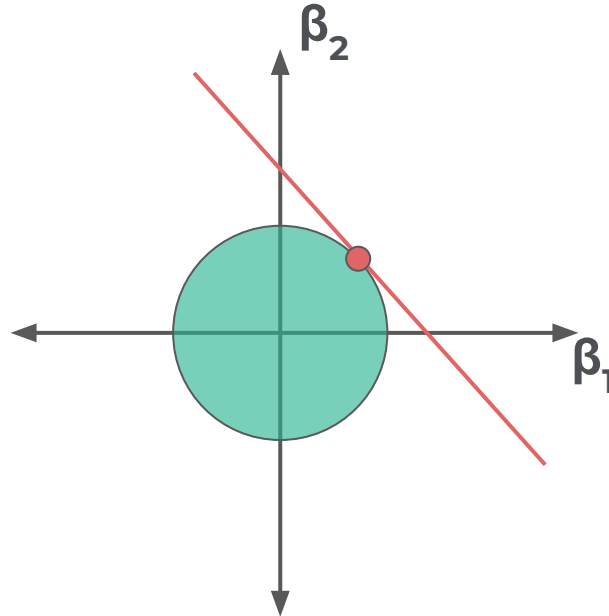
- But were subject to the penalty!





# Elastic Net

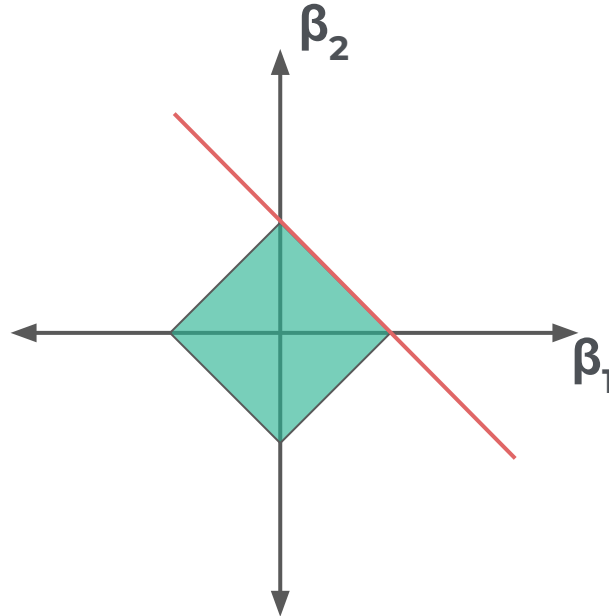
- Penalty for Ridge:  $\beta_1^2 + \beta_2^2 \leq s$





# Elastic Net

- Penalty for Lasso:  $|\beta_1| + |\beta_2| \leq s$





# Elastic Net

- Lasso:
  - A convex object that lies tangent to the boundary, is likely to encounter a corner of a hypercube, for which some components of  $\beta$  are identically zero.



## Elastic Net

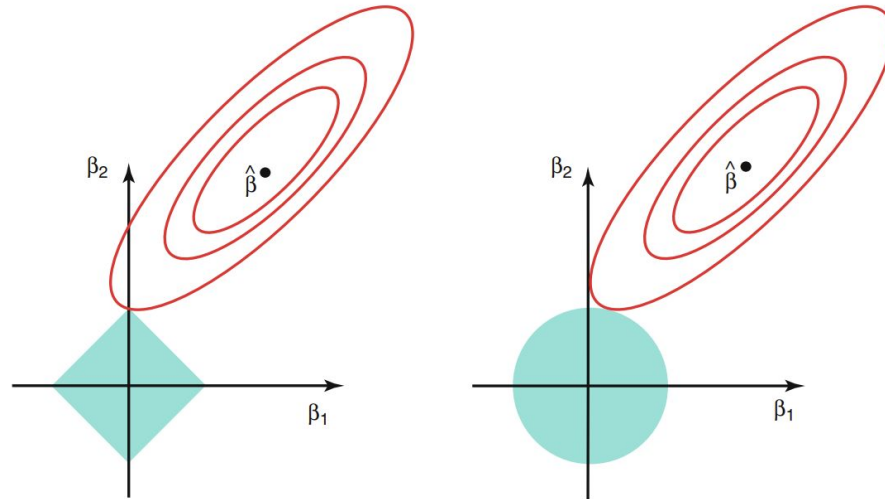
- Ridge: In the case of an n-sphere, the points on the boundary for which some of the components of  $\beta$  are zero are not distinguished from the others and the convex object is no more likely to contact a point at which some components of  $\beta$  are zero than one for which none of them are.





# Elastic Net

- This is why Lasso is more likely to lead to coefficients as zero.
- This diagram is also commonly shown with contour RSS:





## Elastic Net

- Elastic Net seeks to improve on both L1 and L2 Regularization by combining them:

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda_1 \sum_{j=1}^p \beta_j^2 + \lambda_2 \sum_{j=1}^p |\beta_j|$$



## Elastic Net

- Here we seek to minimize RSS and **both** the squared and absolute value terms:

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda_1 \sum_{j=1}^p \beta_j^2 + \lambda_2 \sum_{j=1}^p |\beta_j|$$



# Elastic Net

- Notice there are **two** distinct lambda values for each penalty:

$$\text{Error} = \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda_1 \sum_{j=1}^p \beta_j^2 + \lambda_2 \sum_{j=1}^p |\beta_j|$$



## Elastic Net

- We can alternatively express this as a ratio between L1 and L2:

$$\frac{\sum_{i=1}^n (y_i - x_i^J \hat{\beta})^2}{2n} + \lambda \left( \frac{1 - \alpha}{2} \sum_{j=1}^m \hat{\beta}_j^2 + \alpha \sum_{j=1}^m |\hat{\beta}_j| \right)$$



## Elastic Net

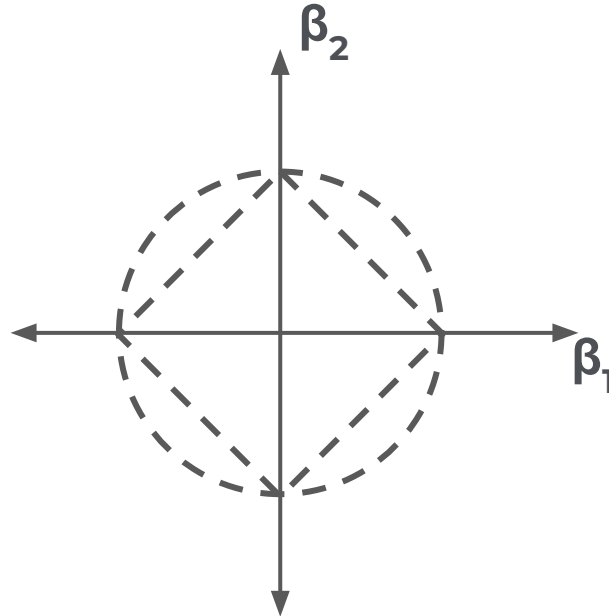
- We can also use simplified notation:

$$\hat{\beta} \equiv \underset{\beta}{\operatorname{argmin}}(\|y - X\beta\|^2 + \lambda_2 \|\beta\|^2 + \lambda_1 \|\beta\|_1)$$



# Elastic Net

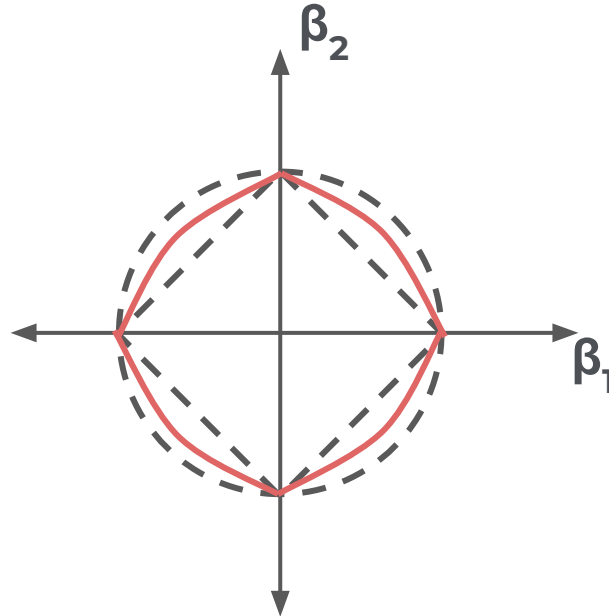
- Elastic Net Penalty Region:





# Elastic Net

- Elastic Net Penalty Region:







# Elastic Net

- Let's explore how to perform Elastic Net with Python and Scikit-learn!

$$\frac{\sum_{i=1}^n (y_i - x_i^J \hat{\beta})^2}{2n} + \lambda \left( \frac{1 - \alpha}{2} \sum_{j=1}^m \hat{\beta}_j^2 + \alpha \sum_{j=1}^m |\hat{\beta}_j| \right)$$